eXpress: Guided Path Exploration for Regression Test Generation

Kunal Taneja[1], Tao Xie[1], Nikolai Tillmann[2], Jonathan de Halleux[2], Wolfram Schulte[2]

[1]Department of Computer Science, North Carolina State University, Raleigh, NC, 27695, USA
[2]Microsoft Research, One Microsoft Way, Redmond, WA, 98074, USA
[1]{ktaneja, txie}@ncsu.edu, [2]{nikolait, jhalleux, schulte}@microsoft.com

ABSTRACT

Regression test generation aims at generating a test suite that can detect behavioral differences between the original and the new versions of a program. Regression test generation can be automated by using Dynamic Symbolic Execution (DSE), a state-of-the-art test generation technique. DSE explores all feasible paths in the program but exploration of all these paths can often be expensive. However, if our aim is to detect behavioral differences between two versions of a program, we do not need to explore all these paths in the program, since not all these paths are relevant for detecting behavioral differences. In this paper, we propose an approach on guided path exploration that avoids exploring irrelevant paths in terms of detecting behavioral differences. Hence, behavioral differences are more likely to be detected earlier in path exploration. In addition, our approach leverages the existing test suite (if available) for the original version to efficiently execute the changed regions of the program and infect program states. Experimental results on 67 versions (in total) of four programs show that our approach requires about 36% fewer amount of time (on average) to detect behavioral differences than exploration without using our approach. In addition, our approach detects 6 behavioral difference that could not be detected by exploration without using our approach (within a time bound). Furthermore, our approach requires 67% fewer amount of time to find behavioral differences by exploiting an existing test suite than exploration without using the test suite.

1. INTRODUCTION

Regression test generation aims at generating a test suite that can detect behavioral differences between the original and the new versions of a program. A behavioral difference between two versions of a program can be reflected by the difference between the observable outputs produced by the execution of the same test (referred to as a difference-exposing test) on the two versions. Developers can inspect these behavioral differences to determine whether they are intended or unintended (i.e., regression faults).

Regression test generation can be automated by using Dynamic Symbolic Execution (DSE) [12, 28, 5], a state-of-the-art test generation technique, to generate a test suite achieving high structural

coverage. DSE explores paths in a program to achieve high structural coverage, and exploration of all these paths can often be expensive. However, if our aim is to detect behavioral differences between two versions of a program, we do not need to explore all these paths in the program since not all these paths are relevant for detecting behavioral differences.

To formally investigate irrelevant paths for exposing behavioral differences, we adopt the Propagation, Infection, and Execution (PIE) model [33] of error propagation. According to the PIE model, a fault can be detected by a test if a faulty statement is executed (E), the execution of the faulty statement infects the state (I), and the infected state (i.e., error) propagates to an observable output (P). A change in the new version of a program can be treated as a fault and then the PIE model is applicable for effect propagation of the change. Our key insight is that many paths in a program often cannot help in satisfying any of the conditions P, I, or E of the PIE model.

In this paper, we present an approach[1] eXpress and its implementation that uses DSE to detect behavioral differences based on the notion of the PIE model.

Our approach first determines all the branches (in the program under test) that cannot help in achieving any of the conditions E, I, and P of the PIE model in terms of the changes in the program. To make test generation efficient, we develop a new search strategy for DSE to avoid exploring these irrelevant branches (including which can lead to an irrelevant path[2]). In particular, our approach guides DSE to avoid from flipping branching nodes[3], which on flipping execute some irrelevant branch.

In addition, our approach can exploit the existing test suite (if available) for the original version by seeding the tests in the test suite to the program exploration. Our technique of seeding the exploration with the existing test suite efficiently augments an existing test suite so that various changed parts of the program (that are not covered by the existing test suite) are covered by the augmented test suite.

This paper makes the following major contributions:

Problem Definition. We formally define the problem of regression

---

---

[1]An earlier version [30] of this work is described in a four-page paper that appears in the NIER track of ICSE 2009. This work significantly extends the previous work in the following major ways. First, in this paper, we develop techniques for efficiently finding irrelevant branches that cannot execute any change. Second, we develop techniques for exploiting the existing test suite for efficiently generating regression tests. Third, we automate our approach by developing a tool. Fourth, we conduct extensive experiments to evaluate our approach.

[2]An irrelevant path is a path that cannot help in achieving P, I, and E of the PIE model.

[3]A branching node in the execution tree of a program is an instance of a conditional statement in the source code. A branching node consists of two sides (or more than two sides for a `switch` statement): the true branch and the false branch. Flipping a branching node is flipping the execution of the program from the true (or false) branch to the false (or true) branch. Flipping a branching node for a switch statement is flipping the execution of the current branch to another unexplored branch.

test generation.

Path Exploration for Regression Test Generation. We propose an approach called `eXpress` that uses DSE for efficient generation of regression tests. To the best of our knowledge, ours is the first approach that guides path exploration on the new version specifically for regression test generation.

Incremental Exploration. We develop a technique for exploiting an existing test suite, so that path exploration focuses on covering the changes rather than starting from scratch.

Implementation. We have implemented our `eXpress` approach in a tool as an extension for Pex [31], an automated structural testing tool for .NET developed at Microsoft Research. Pex has been previously used internally at Microsoft to test core components of the .NET architecture and has found serious faults [31]. The current Pex has been downloaded tens of thousands of times in industry.

Evaluation. We have conducted experiments on 67 versions (in total) of four programs. Experimental results show that our approach requires about 36% fewer amount of time (on average) to detect behavioral differences than exploration without using our approach. In addition, our approach detects 6 behavioral difference that could not be detected by exploration without using our approach (within a time bound). Furthermore, our approach requires 67% fewer amount of time to find behavioral differences by exploiting an existing test suite than exploration without using the test suite.

2. DYNAMIC SYMBOLIC EXECUTION

In our approach, we use Pex [31] as an engine for Dynamic Symbolic Execution (DSE). Pex starts program exploration with some default inputs. Pex then collects constraints on program inputs from the predicates at the branching statements executed in the program. We refer to these constraints at branching statements as branch conditions. The conjunction of all branch conditions in the path followed during execution of an input is referred to as a path constraint. Pex keeps track of the previous executions to build a dynamic execution tree. Pex, in the next run[4], chooses one of the unexplored branches in the execution tree (dynamically built thus far). Pex flips the chosen branching node to generate a new input that follows a new execution path. Pex uses various heuristics [36] for choosing a branching node (to flip next) using various search strategies with an objective of achieving high code coverage fast. We next present definitions of some terms that we use in the rest of this paper.

Discovered node. We refer to all the branching nodes that are explored in the current DSE run but were not explored in previous runs as discovered branching nodes (in short as discovered nodes).

Instance of a branching node. A branching node c_i in a Control Flow Graph (CFG) of a program can be present multiple times in the dynamic execution tree (of the program) due to loops in the program. We refer to these multiple branching nodes in the tree as instances of c_i.

3. EXAMPLE

In this section, we illustrate different components of the `eXpress` approach with an example. `eXpress` takes as input two versions of a program and produces as output a regression test suite, with the objective of detecting behavioral differences (if any exist) between the two versions of program under test. Although `eXpress` analyzes assembly code of C# programs, in this section, we illustrate the `eXpress` approach using program source code.

```
   static public int TestMe(char[] c){
1     int state = 0;
2     if(c == null || c.Length == 0)
3       return -1;
4     for(int i=0; i< c.Length; i++){
5         if(c[i] == "[")
6           state =1;
7         else if(state == 1 && c[i] == "{")
8           state =2;
9         else if(state == 2 &&   c[i] == "<")
10          state =3;
11        else if(state == 3 && c[i] == "*"){
12          state =4;
13          if(c.Length==15) break;//Added in new version
          }
14        if(c[i]==' ')
15          return;
16        if(!(c[i] >= 'a' && c[i] <= 'z'){
17            state=-1; return;
          }
      }
    }
18    if(c[15] == '}')
19      return state;
20    return -1;
   }
```

Figure 1: An example program

Consider the example in Figure 1. Suppose that the statements at Line 13 of `TestMe` are added in the new version. We next describe our approach using the example.

Difference Finder. The Difference Finder component compares the original and the new versions of the program under test to find differences between each corresponding method of the two program versions. For the program in Figure 1, Difference Finder detects that the statements at Line 13 are added in the new version.

Graph Builder. The Graph Builder component then builds a Control-Flow Graph (CFG) of the new version of the program under test and marks the changed vertices in the graph. Figure 3 shows the CFG of the program in Figure 1. The labels of vertices in the CFG denote the corresponding line numbers in Figure 1. The black vertex denotes the newly added statements at Line 13. The gray vertices denote the branching nodes (for the branching statements in the program), while the white vertices denote the other statements in the program.

Graph Traverser. The Graph Traverser component traverses the CFG to find each branch[5] b in a program such that if b is taken, the program execution cannot help in finding behavioral differences between original and the new program versions. These branches are used by the Dynamic Test Generator component to efficiently find behavioral differences between the original and new program version. In particular, the Graph Traverser finds two categories of branches.

- **Category B_{I+E}.** On traversing the CFG in Figure 3, the Graph Traverser detects that taking the branches $< 2, 3 >$, $< 4, 18 >$, $< 16, 17 >$, $< 18, 19 >$, and $< 18, 20 >$ (dotted edges in Figure 3), where $< x, y >$ represents a branch from source vertex x to destination vertex y, the program execution cannot reach the black vertex. Since, after taking these five branches, the execution cannot reach the changed statements at Line 13, the execution of these branches cannot help in executing the changed statements or infecting the program state.

- **Category B_P.** In addition, the Graph Traverser component finds out that among the source vertices of the five branches in Category B_{I+E}, there is no path between the black vertex and the vertex 3. Hence, a state infection after the execution of the black vertex cannot propagate through the branch $< 2, 3 >$.
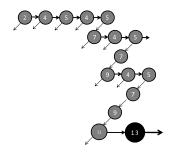

Figure 2: Part of execution tree of the program for c = {"[", "{", "<", "*" }.

Dynamic Test Generator. To cover the changed statements at Line 13 and infect the program state, DSE needs at least 6 DSE runs (starting from an empty input array c). However, the number of runs depends on the choice of branches that DSE flips in each run. In each run, DSE has the choice of flipping 8 branches in the program. For the program in Figure 1, Pex takes 441 DSE runs to cover the true branch of statement at Line 13 (to infect the program state). The number of runs can be much more if the number of branching statements in the program increases. To reduce the search space of DSE, the Dynamic Test Generator component adopts the PIE model [33] of error propagation described in Section 1. In particular, the Dynamic Test Generator prunes various branches from the search space of DSE so that DSE has fewer branches to flip in each run. In particular, the Dynamic Test Generator prunes branches in category B_{E+I} until the black vertex is executed and the program state is infected after its execution. Once the program state is infected, the Dynamic Test Generator prunes out the branches in category B_P to efficiently propagate the state infection to an observable output.

Incremental Exploration. Our approach can reuse an existing test suite for the original version so that changed parts of the program can be explored efficiently due to which test generation is likely to find behavior differences earlier in path exploration. Assume that there is an existing test suite covering all the statements in the old version of the program in Figure 1. Suppose that the test suite has a test input I = {"[", "{", "<", "*" }. The input covers all the statements in the new version of `TestMe` except the newly added `break` statement at Line 13. If we start the program exploration from a scratch (i.e., with default inputs), Pex takes 441 runs to cover the `break` statement at Line 13. However, we can reuse the existing test suite for exploration to cover the new statement efficiently. Our approach executes the test suite to build an execution tree for the tests in the test suite. Our approach then starts program exploration using the dynamic execution tree built by executing the existing test suite instead of starting from an empty tree. Some branches in the tree may take many runs for Pex to discover if starting from an empty tree. Figure 2 shows part of the execution tree for the input I. A gray edge in the tree indicates the false side of a branching node while a black (horizontal) edge indicates the true side. To generate an input for the next DSE runs, Pex flips a branching node b in the tree whose the other side has not yet been explored and generates an input so that program execution takes the unexplored branch of b. Pex chooses such branching node for flipping using various heuristics for covering changed regions of the program. It is likely that Pex chooses the branching node 13 (colored black), which on execution covers the `break` statement at Line 13. When starting the program exploration from scratch, Pex takes 420 runs before it reaches (discovers) the black branching node in Figure 2. Using our approach of seeding the tests from the existing test suite, Pex takes 39 runs to flip the branching node and cover the break statement at Line 13.
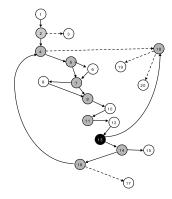

Figure 3: The Control-Flow Graph for the program in Figure 1

4. PROBLEM DEFINITION

In this section, we present formal definitions of terms that we use throughout this paper. In addition, we formally define the regression test generation problem that we address in this paper.

We use concrete-state representation from our previous work [34] to represent the state of an object. The formal definitions of Object State and Object State Equivalence can be found in our previous work [34].

DEFINITION 4.1. **Program Execution Point.** Let $S = (s_1, s_2, ..., s_n)$ be the sequence of statements executed during a program execution $V(I)$ (i.e., when program V is executed with a list of argument values I). Since a statement s can be executed multiple times, multiple instances of s can be present in S. A program execution point (in short an execution point) for a program execution $V(I)$ is a statement $s_i \in S$ such that statements $(s_1, s_2..., s_i)$ have already been executed but statements $s_{i+1}, s_{i+1}, ..s_n$ are yet to be executed.

DEFINITION 4.2. **Program State.** Let $M = \{m_1, m_2, m_3, ..., m_n\}$ be the set of methods in the program call stack at an execution point e. Let $R = \{o_1, o_2, o_3, ..., o_l\}$ be the set of objects such that o_i is the receiver object of m_i, $A = \{a_1, a_2, ..., a_m\}$ be the set of all the arguments (primitive or non primitive) to methods in M, and $L = \{l_1, l_2, l_3, ..., l_x\}$ be the set of all the in-scope local and global variables (primitive or non-primitive) at the statements $P = \{p_1, p_2, p_3, ..., p_y, s_e\}$, where p_i is the statement at which method m_i is invoked and the s_e is the statement corresponding to the execution point e. The program state at an execution point e is the set of object states $S_R \cup S_A \cup S_L$, where $S_R = \{S_{o1}, S_{o2}, S_{o3}..., S_{ol}\}$ (S_{oi} is the object state for object o_i at e), $S_A = \{S_{a1}, S_{a2}, S_{a3}..., S_{an}\}$ (S_{ai} is the object state for a_i if a_i is non-primitive otherwise S_{ai} is the value of a_i), and $S_L = \{S_{l1}, S_{l2}, S_{l3}..., S_{lx}\}$ (S_{li} is the object state for l_i at s_e if l_i is non-primitive otherwise S_{li} is the value of l_i).

Let a program version V be modified to incorporate a set of changes Δ. We assume that signatures of methods in V are not changed and no new fields are added, modified, or deleted in V. Δ also includes the new statements added to V and statements deleted from V. The program version V' is the resulting program version after Δ is applied to the program version V. A method M in V corresponds to a method M' (denoted as $M \leftrightarrow M'$) in V' if M and M' have the same signatures. Assume that some dummy statements are added to the two program versions for matching added/deleted statements between the two versions (as done by Santelices et al. [26]). These dummy statements make sure that a statement s in V can be mapped to a corresponding statement s' in V'.

DEFINITION 4.3. Execution Index [37]. An execution Index $E_{V(I)}$ of a program execution $V(I)$ is the function of execution points in $V(I)$ that satisfies the following property: for any two execution points $x \in V(I)$ and $y \in V(I)$: if $x \neq y \Rightarrow E_{V(I)}(x) \neq E_{V(I)}(y)$

DEFINITION 4.4. Corresponding Execution Points [37]. An execution point e in the execution $V(I)$ corresponds to an execution point e' (denoted as $e \leftrightarrow e'$) in the execution $V'(I)$ (V' is a modified version of V) when $E_{V(I)}(e) = E_{V'(I)}(e')$. For an execution point e in the execution $V(I)$, there may or may not be a corresponding execution point e' in the execution $V'(I)$.

DEFINITION 4.5. Parameterized Unit Test (PUT). In our formal model, we assume that a program is invoked through a Parametrized Unit Test (PUT) [32]. A PUT $P_V = \langle A, S, \lambda \rangle$ is a method with a list of arguments A and a list of statements $S = (s_1, s_2, ..s_n)$ such that $s_i \in S$ is a statement other than an assertion statement and λ is a list of assertion statements.

DEFINITION 4.6. Corresponding PUTs. In our formal model, we assume that a program is invoked through a Parametrized Unit Test (PUT) [32]. A PUT $P_V = \langle A, S, \lambda \rangle$ in a program version V, corresponds to a PUT $P_{V'} = \langle A', S', \lambda' \rangle$ (denoted as $P_V \leftrightarrow P'_V$), in program version V' if V and V' are different versions of a same program, $A = A'$, $S = S'$, and $\lambda = \lambda'$.

DEFINITION 4.7. Infected Program State. The program state of a program version V' (invoked from PUT $P_{V'}$) is infected with respect to a program version V (invoked by PUT P_V: $P_V \leftrightarrow P'_V$) at an execution point $e' \in V'(I)$, if $\exists e \in V(I) : e \leftrightarrow e'$, program state of V' at e' is S' and program state of V at $e \in V(I)$ is S, and S is not equivalent to S' (definition of equivalent states can be found in our previous work [34]).

DEFINITION 4.8. Behavioral Difference. Two corresponding PUTs P_V and $P_{V'}$ are invoked with a list of argument values (input) $I = (a_1, a_2, ..., a_k)$. There is a behavioral difference between V and V' if there exists an input I such that when P_V and $P_{V'}$ are invoked with input I, $\exists \lambda_i \in \lambda$ such that λ_i passes for P_V but fails for P'_V or λ_i fails for P_V but passes for P'_V. I is referred to as an input detecting behavioral difference.

DEFINITION 4.9. Problem Definition. Regression test generation for a PUT $P_{V'}$ in program version V' is the problem of generating an input I for $P_{V'}$ that detects behavioral difference between V and V'. This paper focuses on efficiently generating inputs detecting behavioral differences.

5. APPROACH

eXpress takes as input the assembly code of two versions $v1$ (original) and $v2$ (new) of the program under test. In addition, eXpress takes as input, the name and signature of a parameterized unit test (PUT); such a PUT serves as a test driver for path exploration. eXpress consists of five main components.

The Difference Finder component finds the set of differences between each of the corresponding method pairs in the two versions of the program. The Graph Builder component builds an inter-procedural graph G with the input PUT as the starting method. The Graph Traverser component traverses the graph to find all the branches B_{E+I} that need not be explored for executing the changed regions and infect the program state and the branches B_P that cannot help in propagating a state infection to an observable output. The Instrumenter component instruments both versions of the program under test so that program states can be compared (to determine state infection). The Dynamic Test Generator generates tests

for the input PUT, pruning branches B_{E+I} from exploration before a test is generated that infects the program state. In addition, Dynamic Test Generator prunes branches B_P after the program state is infected until a behavioral difference is found.

When an existing test suite is available for the original version, eXpress conducts incremental exploration that exploits the test suite for generating tests for the new version. We next discuss in detail the main components in eXpress and its incremental exploration.

5.1 Difference Finder

The Difference Finder component takes the two versions $v1$ and $v2$ as input and analyzes the two versions to find pairs $\langle M_{i1}, M_{i2} \rangle$ of corresponding methods in $v1$ and $v2$, where M_{i1} is a method in $v1$ and M_{i2} is a method in $v2$. A method M is defined as a triple $\langle FQN, Sig, I \rangle$, where FQN is the fully qualified name[6] of the method, Sig is the signature[7] of the method, and I is the list of assembly instructions in the method body. Two methods $\langle M_{i1}, M_{i2} \rangle$ form a corresponding pair if the two methods M_{i1} and M_{i2} have the same FQN and Sig. Currently our approach considers as different methods the methods that have undergone Rename Method or Change Method Signature refactoring. A refactoring detection tool [7] can be used to find such corresponding methods. For each pair $\langle M_{i1}, M_{i2} \rangle$ of corresponding methods, the Difference Finder finds a set of differences Δ_i between the list of instructions $I_{M_{i1}}$ and $I_{M_{i2}}$ in the body of Methods M_{i1} and M_{i2}, respectively. Δ_i is a set of instructions such that each instruction ι in Δ_i is an instruction in $I_{M_{i2}}$ (or in $I_{M_{i1}}$ for a deleted instruction), and ι is added, modified, or deleted from list $I_{M_{i1}}$ to form $I_{M_{i2}}$.

In its other components, eXpress analyzes Version $v2$ of the program while using the differences obtained from the Difference Finder component to efficiently generate regression tests. Note that Version $v1$ can also be used instead of $v2$ in path exploration.

5.2 Graph Builder

The Graph Builder component makes an inter-procedural control flow graph (CFG) of the program version v_2. The Graph Builder component starts the construction of the inter-procedural CFG from the Parametrized Unit Test (PUT) τ provided as input. The inter-procedural CFG is used by the Graph Traverser component to find branches (in the graph) via which the execution cannot reach any vertex containing a changed instruction in the graph. Since a moderate-size program can contain millions of method calls (including those in its dependent libraries), often the construction of inter-procedural graph is not scalable to real-world programs. Hence, we build a minimal inter-procedural CFG for which our purpose of finding branches that cannot reach some changed region in the program can be served. The pseudo code for building the inter-procedural CFG is shown in Algorithm 1. Initially, the algorithm InterProceduralCFG is invoked with the argument as the PUT τ. The algorithm first constructs an intra-procedural CFG g for method τ. For each method invocation vertex[8] (invoking Method c) in g, the algorithm InterProceduralCFG is invoked recursively with the invoked method c as the argument (Line 22 of Algorithm 1), after adding c to the call stack (Line 21). After the control returns from the recursive call, the method c is removed from the call stack (Line 23) and added to the set of visited meth-

---

[6] The fully qualified name of a method m is a combination of the method's name, the name of the class c declaring m, and the name of the namespace containing c.

[7] Signature of a method m is the combination of parameter types of m and the return type of m.

[8] A method invocation vertex is a vertex representing a call instruction.

Algorithm 1 $InterProceduralCFG(\tau)$

Input: A test method τ.
Output: The inter-procedural Control Flow Graph (CFG) of the program under test.

```
1:  Graph g ← GenerateIntraProceduralCFG(τ)
2:  for all Vertex v ∈ g.Vertices do
3:      if v.Instruction = MethodInvocation then
4:          c ← getMethod(v.Instruction)
5:          if c ∈ MethodCallStack then
6:              goto Line 2 //To avoid loops
7:          end if
8:          if c ∈ ReachableToChangedRegion then
9:              g ← GraphUnion(ChangedMethod, g, v)
10:             goto Line 2
11:         end if
12:         if c ∈ Visited then
13:             goto Line 2
14:         end if
15:         if c ∈ ChangedMethods then
16:             ChangedMethod ← c
17:             for all Method m ∈ MethodCallStack do
18:                 ReachableToChangedRegion.Add(m)
19:             end for
20:         end if
21:         MethodCallStack.Add(c)
22:         cg ← InterProceduralCFG(c)
23:         MethodCallStack.Remove(c)
24:         Visited.Add(c)
25:         g ← GraphUnion(cg, g, v)
26:     end if
27: end for
28: return g
```

ods (Line 24). The inter-procedural graph cg (with c as an entry method) resulting from the recursive call at Line 22 is merged with the graph g (Line 25). The algorithm `InterProceduralCFG` is not invoked recursively with c as argument in the following situations:

c **is in call stack.** If c is already in the call stack, `InterProcedural CFG` is not recursively invoked with c as argument (Lines 5-6). This technique ensures that our approach is not stuck in a loop in method invocations. For example, if method A invokes method B, and method B invokes A, then the construction of the inter-procedural graph stops after A is encountered the second time.

c **is already visited.** If c is already visited, `InterProcedural CFG` is not recursively invoked with c as argument (Line 23). This technique ensures that we do not build the same subgraph again.

c **is in** `ReachableToChangedRegion`. The set `ReachableToChanged Region` is populated whenever a changed method[9] is encountered. In particular, if a changed method is encountered, the methods currently in the call stack are added to the set `ReachableToChanged Region` (Lines 17-19). If c is in `ReachableToChangedRegion`, `InterProceduralCFG` is not recursively invoked with c as argument, while merging CFG of some changed method with g (Line 8-11). Note that if a method m can reach a changed method cm, all the branches in this method m may not be able to reach cm (for example, due to $return$ statements). Branches that cannot reach any changed region (irrelevant branches) are found by the Graph Traverser component.

If a branching node b can reach a changed region in Inter-Procedural CFG $g1$ built without using the preceding optimization, the branching node b can reach a changed region (maybe a different one) in graph $g2$ built using the preceding optimizations. Since our aim of building the intra-procedural CFG is to find irrelevant branches, those in the graph via which the execution cannot reach any changed region, the preceding three optimizations help achieve the aim while reducing the cost of building the inter-procedural CFG. In addition, the size of the inter-procedural CFG is reduced resulting in reduction in the cost of finding irrelevant branches.

5.3 Graph Traverser

The Graph Traverser component takes as input the inter-

---

[9] A changed method M_i is a method for which the set $\Delta_i \neq \phi$.

procedural CFG $g < V, E >$ constructed by the Graph Builder and a set C of changed vertices in the CFG g (found by Difference Finder). The Graph Traverser traverses the graph to find a set of branches B_{E+I} via which the execution cannot reach any of the branches in V and a set of branches B_P via which a state infection cannot propagate to any observable output. A branch b in CFG g is an edge $e = < v_i, v_j >: e \in E$, $v_i \in V$ with an outgoing degree of more that one. The vertex v_i is referred to as a branching node. We next describe the sets B_{E+I} and B_P.

Let $V = \{v_1, v_2, .., v_l\}$ be the set of all vertices in CFG $g < V, E >$ such that $v_i \in V$ and $v_i.degree > 1$. Let $E_i = \{e_{i1}, e_{i2}, ..., e_{im}\}$ be the set of outgoing edges (branches) from v_i. Let C be the set of changed vertices in g, $\rho(v_i, v_j, e_{ij})$ denotes a path between a source vertex v_i to a destination vertex v_j that takes the branch e_{ij} (if $\rho(v_i, v_j, e_{ij}) = \phi$, there is no such path from v_i to v_j), $\rho(v_i, v_j)$ denotes a path from a source vertex v_i to a destination vertex v_j(if $\rho(v_i, v_j) = \phi$, there is no such path from v_i to v_j).

Branches B_{E+I}. $B_{E+I} \subseteq E$ is a set of branches such that $\forall e_{ij} = < v_i, v_j > \in B_{E+I} \land \forall c_k \in C : \rho(v_i, c_k, e_{ij}) = \phi$. Since $\rho(v_i, c_k, e_{ij}) = \phi$, after taking a branch $e_{ij} \in B_{E+I}$, the program execution cannot reach a changed region. Hence, the program state cannot be infected.

Branches B_P. $B_P \subseteq E$ is a set of branches such that $\forall e_{ij} = < v_i, v_j > \in B_P \land \forall c_k \in C : \rho(v_i, c_k, e_{ij}) = \phi \land \rho(c_k, v_i) = \phi$. Since $\rho(c_k, v_i) = \phi$, a state infection after a changed vertex c_k cannot reach v_i. Hence, the state infection cannot propagate through e_{ij}.

The Graph Traversal component uses Depth First Search on the Graph g to efficiently find the sets B_{E+I} and B_P.

5.4 Instrumenter

The Instrumenter component uses Sets Δ_i (differences between method M_{i1} and M_{i2}) produced by the Difference Finder. For each changed method pair $< M_{i1}, M_{i2} >$ for which $\Delta_i \neq \phi$, the Instrumenter component finds a region δ_i containing all the changed instructions in the program. δ_i is a minimal list of continuous instructions such that all the changed instructions in the method M_i are in the region δ_i. Hence, there can be a maximum of one changed region in one method. At the end of each changed region δ_i, the Instrumenter component inserts instructions to save the program state. In particular, the Instrumenter inserts the corresponding instructions for `PexStore` statements for each variable (and field) defined in the changed region. The `PexStore` statement for a variable x results in an assertion statement `PexAssert.IsTrue( "uniqueName", x == currentX)` in the generated test, where `currentX` is the value of x in the new version. The Dynamic Test Generator component generates tests for the new version $v2$. After the Dynamic Test Generator component generates a test that executes a changed region (note that in $v1$, we do not have a changed region), the component executes the test on Version $v1$ to compare program states after the execution of the changed region with the ones captured in the execution of Version $v2$. The instrumentation enables us to perform only one instance of DSE on the new version instead of performing two instances of DSE: one on the original and the other on the new program version. Performing two instances of DSE can be technically challenging since we have to perform the two DSE instances in a controlled manner such that both versions are executed with the same input and the execution trace is monitored for both the versions by a common exploration strategy to decide which branching node to flip next in the two versions.

5.5 Dynamic Test Generator

The Dynamic Test Generator component performs Dynamic Symbolic Execution (DSE) [6, 17, 12, 28, 5] to generate regression tests for the two given versions of a program. DSE iteratively generates test inputs to cover various feasible paths in the program under test (the new version in our approach). In particular, DSE flips some branching node discovered in previous executions to generate a test input for covering a new path. The branching node to be flipped is decided by a search strategy such as depth-first search. The exploration is quite expensive since there are an exponential number of paths with respect to the number of branches in a program. However, the execution of many branches often cannot help in detecting behavioral differences. In other words, covering these branches does not help in satisfying any of the condition E or I in the PIE model described in Section 1. Therefore, we do not flip such branching nodes in our new search strategy for generating test inputs that detect behavioral differences between the two given versions of a program. Recall that, we refer to such branches as irrelevant branches. These branches are found using the Graph Traverser component. We next describe the two categories of paths that our approach avoids exploring. We then describe our incremental exploration technique when an existing test suite is available for the original version.

5.5.1 Path Pruning

Our approach avoids exploring the following categories of paths:
Category E+I. eXpress does not explore these paths until a test is generated that infects the program state. Our approach avoids exploring all the branches in B_{E+I}. In other words, if a branching node v_i is in the built dynamic execution tree[10], such that the branch $e_{ij} \in B_{E+I}$ is not explored yet, our approach does not flip v_i to make program execution take branch e_{ij}. As an effect, our approach helps reduce the number of paths to be explored by avoiding the paths $P \in g : e_{ij} \in P$ until a test is generated that infects the program state.
Category P. eXpress does not explore these paths after a test is generated that infects the program state. Our approach avoids exploring all the branches $e_{ij} \in B_p$. In other words, if a branching node v_i is in the built dynamic execution tree such that one of its branch e_{ij} is not explored yet, our approach does not flip v_i to make program execution take branch e_{ij}. As an effect, our approach helps reduce the number of paths to be explored by avoiding the paths $P \in g : e_{ij} \in P$. These paths cannot help in propagating a state infection to an observable output.

5.5.2 Incremental Exploration

A regression test suite achieving high code coverage may be available along with the original version of a program. This test suite may be manually written or generated by an automated test generation tool for the original version. However, the existing test suite might not be able to cover all the changed regions of the new version of the program. Our approach can reuses the existing test suite so that changed regions of the program can be executed efficiently due to which test generation is likely to find behavior differences earlier in path exploration. Our approach executes the existing test suite to build an execution tree for the tests in the test suite. Our approach then starts the program exploration using the dynamic execution tree instead of starting from an empty tree. Our approach of seeding test inputs can help efficiently cover the changed regions of the program with two major reasons:
Discovery of hard-to-discover branching nodes. By seeding the

existing test suite for Pex to start exploration with, our approach executes the test suite to build an execution tree of the program. Some of the branching nodes in the built execution tree may take a large number of DSE runs (without seeding any tests) to get discovered. Flipping some of these discovered branching nodes nearer in CFG to the changed parts of the program has more likelihood of covering the changed regions of the program [4]. Although, our approach currently does not specifically prioritize flipping of branching nodes near the changed regions, our approach can help these branching nodes to get discovered (by executing the existing test suite), which might take large number of DSE runs as shown in the example in Section 3.
Priority of DSE to cover not-covered regions of the program. DSE techniques employ branch prioritization so that a high coverage can be achieved faster due to which DSE techniques choose a branch from the execution tree (built thus far) that have a high likelihood of covering changed regions (that are not covered by existing test suite for the original version). By seeding the existing test suite to program exploration, the DSE techniques do not waste time on covering the regions of the program already covered by the existing test suite. Instead, the DSE techniques give high priority to branching nodes that can cover not-covered regions of the program, which include the changed parts. Hence, the changed parts are likely to be covered earlier in path exploration.

6. EXPERIMENTS

We conducted experiments on four programs and their 67 versions (in total) collected from three different sources. In our experiments, we try to address the following research questions:
RQ1. How many fewer DSE runs does Pex require to execute the changed regions between the two versions of a program with the assistance of eXpress?
RQ2. How many fewer DSE runs does Pex require to infect the program states with the assistance of eXpress?
RQ3. How many fewer DSE runs and how much fewer amount of time does Pex require to find behavioral differences with the assistance of eXpress?
RQ4. How many fewer DSE runs and how much fewer amount of time does Pex require to find behavior differences when the program exploration is seeded with the existing test suite?

6.1 Subjects

To answer the research questions, we conducted experiments on four subjects. Table 1 shows the details about the subjects. Column 1 shows the subject name. Column 2 shows the number of classes in the subject. Column 3 shows the number of classes that are covered by tests generated in our experiments. Column 4 shows the number of versions (not including the original version) used in our experiments. Column 5 shows the number of lines of code in the subject.

replace and siena are programs available from the Subject Infrastructure Repository (SIR) [8]. replace and siena are written in C and $Java$, respectively. replace is a text-processing program, while siena is an Internet-scale event notification program. We chose these two subjects (among the others available at the SIR) in our experiments we could convert these subjects into C# using the Java 2 CSharp Translator[11]. We could not convert other subjects available at the SIR (with the exception of tcas) because of extensive use of C or $Java$ library APIs in these subjects. The experimental results on tcas are presented in a previous version

---

[10]A dynamic execution tree is the tree formed from the paths executed in the previous executions. Multiple instances $c_{i1}, c_{i2}, ..., c_{in}$ of a node c_i in CFG can be present in a dynamic execution tree due to loops in a program.

of this work [30] and show similar conclusions as the results from the subjects used in the experiments here. We seeded all the 32 faults available for `replace` at the SIR one by one to generate 32 new versions of `replace`. For `siena`, SIR contains 8 different sequentially released versions of `siena` (versions 1.8 through 1.15). Each version provides enhanced functionalities or corrections with respect to the preceding version. We use these 8 versions in our experiments. In addition to these 8 versions, there are 9 seeded faults available at SIR. We seeded all the 9 faults available at SIR one by one to synthesize 9 new versions of `siena`. In total, we conduct experiments on these 17 versions of `siena`. For `replace`, we use the `main` method as a PUT for generating tests. We capture the concrete value of the string sub at the end of the PUT using `PexStore.ValueForValidation("v", v)` statement. This statement captures the current value of v in a particular run (i.e., an explored path) of DSE. In particular, this statement results in an assertion `Assert.AreEqual(v, cv)` in a generated test, where cv is the concrete value of v in the test during the time of exploration. This assertion is used to find behavioral differences when the tests generated for a new version are executed on the original version.

For `siena`, we use the methods `encode` (for changes that are transitively reachable from `encode`) and `decode` (for changes that are transitively reachable from `decode`) in the class `SENP` as PUTs for generating tests. We capture the return values of these methods using the `PexStore` statement in the PUTs.

The method `encode` requires non-primitive arguments. Existing Pex cannot handle non-primitive argument types effectively but provides support for using factory methods for non-primitive types. Hence, we manually wrote factory methods for the non-primitive types in `SENP`. In particular, we wrote factory methods for classes `SENPPacket`, `Event` and `Filter`. Each factory method invokes a sequence (of length up to three) of the public state-modifying methods in the corresponding class. The parameters for these methods, and the length of the sequence (up to three) are passed as inputs to the factory methods. During exploration, Pex generates concrete values for these inputs to cover various parts of the program under test.

STPG[12] is an open source program hosted by the codeplex website, Microsoft's open source project hosting website. The codeplex website contains snapshots of check-ins in the code repositories for STPG. We collect three different versions of the subject STPG from the three most recent check-ins. We use the `Convert(string path)` method as the PUT for generating tests since `Convert` is the main conversion method that converts a string path data definition to a `PathGeometry` object. We capture the return value of `Convert` using the `PexStore` statement in the PUTs.

`structorian`[13] is an open source binary data viewing and reverse engineering tool. `structorian` is hosted by Google's open source project hosting website. The website also contains snapshots of check-ins in the code repositories for `structorian`. We collected all the versions of snapshots for the classes `StructLexer` and `StructParser`. We chose these classes in our experiments due to three factors. First, these classes have several revisions available in the repository. Second, these classes are of non-trivial size and complexity. Third, these classes have corresponding tests available in the repository. For the classes `StructLexer` and `StructParser`, we generalized one of the available concrete test methods by promoting primitive types to arguments of the test methods. Furthermore, we convert the assertions in the concrete test methods to `PexStore` statements. Foe example if an asser-

[12]http://stringtopathgeometry.codeplex.com/
[13]http://code.google.com/p/structorian/

Table 1: Experimental subjects

Project	Classes	Classes Covered	Versions	LOC
replace	1	1	32	625
STPG	1	1	2	684
siena	6	6	17	1529
structorian	70	8	21	6561

tion `Assert.IsEqual(v, 0)` exist in a concrete test, we convert the assertion to `PexStore.ValueForValidation("v", v)`. We use these generalized test methods as PUTs for our experiments. `structorian` contains a manually written test suite. We use this test suite for seeding the exploration for addressing the question RQ4.

To address questions RQ1-RQ3, we use all the four subjects, while to address question RQ4, we use `structorian` because of two major factors. First, `structorian` has a manually written test suite that can be used to seed the exploration. Second, revisions of `structorian` contains non-trivial changes that cannot be covered by the existing test suite. Hence, our technique of seeding the existing test suite in the program exploration is useful for covering these changes. `replace` contains changes to one statement due to which most of the changes can be covered by the existing test suite. Hence, our Incremental Exploration technique is not beneficial for the version pairs under test. `siena` and `STPG` do not have an existing test suite to use.

6.2 Experimental Setup

For `replace` and `siena`, we conduct regression test generation between the original version and each version $v2$ synthesized from the available faults in the SIR. We use `eXpress` and the default search strategy in Pex [31, 36] to conduct regression test generation. In addition to the versions synthesized by seeding faults, we also conduct regression test generation between each successive versions of `siena` (versions 1.8 through 1.15) available in SIR, using `eXpress` and the default search strategy in Pex [31, 36]. For STPG and `structorian`, we conduct regression test generation between two successive version pairs that we collected.

To address RQ1, we compare the number of runs of DSE required by the default search strategy in Pex (in short as Pex) with the number of runs required by Pex enhanced with `eXpress` (in short as Pex+eXpress) to execute a changed region. To address RQ2, we compare the number of runs required by Pex with the number of runs required by Pex+eXpress to infect the program states. To address RQ3, we compare the number of runs and the amount of time required by Pex with the number of runs required by Pex+eXpress to find behavioral differences between two versions of a program under test. To address RQ4, we compare the number of DSE runs and the amount of time required by Pex (and Pex+eXpress) to find behavioral differences with and without seeding the program exploration (with the existing test suite).

Currently, we have not automated the Instrumenter component of `eXpress`. To find out whether the program state has been infected after the execution of a changed region, we manually add statements to the code under test that send a callback to the Dynamic Test Generator whenever the program state is infected after the execution of a changed region. The rest of the approach is fully automated and is implemented in a tool as an extension[14] to Pex [31]. We developed its components to statically find irrelevant branches as a .NET Reflector[15] AddIn. In future, we plan to automate the Instrumenter component.

To find behavioral differences between two versions, we execute

[14]http://pexase.codeplex.com/
[15]http://www.red-gate.com/products/reflector/

the tests generated for a new version on the original version. Behavioral differences are detected by a test if an assertion in the test fails.

6.3 Experimental Results

Table 2 shows the experimental results. Due to space limit, we provide only the total, average, and median values for the subjects `replace`, `siena`, and `STPG`. The detailed results for experiments on all the versions of these subjects are available on our project web[16]. However, we provide detailed results for `structorian` in this paper.

Column S shows the name of the subject. For `structorian`, the column shows the class name. The class `StructLexer` is denoted by SL and the class `StructParser` is denoted by SP. Column V shows the number of version pairs for which we conducted experiments for the subject. For `structorian`, the column shows the version numbers on which the experiments were conducted. These version numbers are the revision numbers in the google code repository of `structorian`. Column E_{Pex} shows the total number of DSE runs required by Pex for satisfying E. Column E_{Red} shows the average percentage reduction in the number of DSE runs by Pex+eXpress for achieving E. Column M_e shows the median percentage reduction in the number of DSE runs by Pex+eXpress for achieving E. Column I_{Pex} shows the total number of DSE runs required by Pex for satisfying I. Column I_{Red} shows the average percentage reduction in the number of DSE runs by Pex+eXpress for achieving I. Column M_i shows the median percentage reduction in the number of DSE runs by Pex+eXpress for achieving I. Column P_{Pex} shows the total number of DSE runs required by Pex for satisfying P. Column p_{Red} shows the average percentage reduction in the number of DSE runs by Pex+eXpress for achieving P (i.e., finding behavioral differences). Column M_p shows the median percentage reduction in the number of DSE runs by Pex+eXpress for achieving P. Column T_{Ppex} shows the time taken by Pex for satisfying P. Column $T_s + T_d$ shows the time taken $Pex + eXpress$ for satisfying P. This time includes the time taken to statically find irrelevant branches. Column T_{PRed} shows the average percentage reduction in amount of time taken by Pex+eXpress for achieving P.

Table 3 shows the time taken for finding the irrelevant branches and the number of irrelevant branches found. Column S shows the subject. Column T_{static} shows the average time taken by `eXpress` to find irrelevant branches that cannot help in satisfying E of the PIE model. Column B_{E+I} shows the average number of branches found in the set B_{E+I}. In general, irrelevant branches are more if changes are towards the beginning of the PUT since there are likely to be more branches in the program that do not have a path to any changed regions. These branches also include the branches whose branching conditions are not dependent on the inputs of the program and therefore do not correspond to branching conditions during path exploration. Hence, pruning these branches is not helpful in making DSE efficient. Column B_P shows the average number of branches found in the set B_P. Column B_{Tot} shows the total number of branches in the CFG.

Results of replace. For the `replace` subject, among the 32 pairs of versions, the changed regions cannot be executed for 4 of theses version pairs (version pairs 14, 18, 27, and 31) by Pex or by Pex+eXpress in 1000 DSE runs. We do not include these version pairs while calculating the sum of DSE runs for satisfying I and E of the PIE model. For 3 of the version pairs (version pairs 12, 13, and 21), the changes are in the fields due to which there are no benefits of using Pex+eXpress. We exclude these three version pairs

[16]`https://sites.google.com/site/asergrp/ projects/express/`

from the experimental results shown in Table 2, which includes the results of 32 version pairs. For 3 of the version pairs (version pairs 3, 22 and 32), a changed region was executed but the program state is not infected (both by Pex and Pex+eXpress) in a bound of 5 minutes. We do not include these version pairs while calculating the sum of DSE runs for satisfying I of the PIE model. In addition, for 6 of the version pairs, the state infection was not propagated to observable output within a bound of 5 minutes. We do not include these version pairs while calculating the sum of DSE runs for satisfying P of the PIE model.

In total, Pex+eXpress took 49% fewer runs in executing the changes with a maximum of 77.6% for versions pairs 23 and 24. For these version pairs, Pex+eXpress takes 95 DSE runs in contrast to 425 runs taken by Pex to execute the changed locations. For many version pairs, there was no benefit of using Pex+eXpress in terms of satisfying E of the PIE model. As a result, median reduction in the number of runs is 0% for satisfying E of the PIE model.

Pex+eXpress took 46% fewer runs, in infecting the program state, with a maximum of 73.8% for version pair 6. For this version, eXpress takes 83 DSE runs in contrast to 317 runs taken by Pex to infect the program state after the execution of changed regions. We observe that Pex+eXpress took 75% fewer runs (median 49%) and 67% fewer amount of time in finding behavioral differences. This difference is substantially larger than the reduction in runs to achieve I. This phenomenon is due to a large number of exception paths (i.e., paths that lead to an exception statement such as `throw`) in $replace$. As a result, state infections often do not propagated to an assertion violation in the PUT due to exceptions thrown in the program `replace`. `eXpress` helps in pruning these irrelevant paths that lead to an exception.

Results of siena. We observe that the behavioral differences between five of the version pairs of `siena` are found within ten runs by Pex and Pex+eXpress. For these version pairs, there is no reduction in the number of runs . The reason for the preceding phenomenon is that changes in these version pairs are close to the entry vertex in the CFG. Hence, these changes can be covered in a relatively small number of runs. In two of the version pairs, changed regions were not covered by either Pex+eXpress or Pex. An exception is thrown by the program before these changes could be executed. Pex and Pex+eXpress were unable to generate a test input to avoid the exception. Changes between two of the version pairs were refactorings due to which the program state is never infected.

For two of the changes, behavioral differences could not be detected by Pex within a bound of 5 minutes but they were detected by Pex+eXpress. We use 5 minutes for calculating the total values in the column T_{p_pex} and the number of DSE runs performed during five minutes in the column P_{Pex}.

In summary, Pex+eXpress executed the changed region in 42% fewer runs (median 13%), infected the program state in 39.4% (median 13%) fewer runs, and found behavioral differences in 42% fewer runs(median 15%) and 38% fewer amount of time than Pex. In addition, Pex+eXpress detected behavioral differences for two changes that were not detected by Pex within a bound of 5 minutes.

Results of structorian. The six rows with SL in column S of Table 2 show the experimental results for changes in the class `StructLexer`, while the last 10 rows (with SP in column S) show the experimental results on versions of the class `StructParser`. For the versions of `StructLexer`, Pex+eXpress takes 24% fewer runs (median 26.5%) to execute a changed region than Pex. In addition, Pex+eXpress infects the program state in 24% fewer runs(median 26.5%). Neither Pex nor Pex+eXpress were able to find behavioral differences for two versions. Pex+eXpress takes

Table 2: Experimental results

S	V	Execution			Infection			Propagation					
		E_{Pex}	$E_{Red}(\%)$	M_i	I_{Pex}	$I_{Red}(\%)$	M_e	P_{Pex}	$P_{Red}(\%)$	M_p	Tp_{Pex}	$Ts+T_d$	$Tp_{Red}(\%)$
replace	32	1668	49	0	2740	42	34	10312	75	49	711	235	67
siena	17	286	42	13	284	39.4	13	7301	42	15	1011	628	38
STPG	2	341	26	26	378	32		378	32	32	353	255	28
Total	51	2573	49		3402	41		17613	62		1722	863	50
					—structorian—								
SL	2-9	102	26.5	-	102	26.5	-	-	-	-	-	-	-
SL	9-139	102	26.5	-	102	26.5	-	2988	66	-	99	69.3	68
SL	139-150	102	26.5		102	26.5	-	761	69	-	26	7.5	71
SL	150-169	53	13.2	-	53	13.2	-	299	52	-	7.4	3.9	47
SL	169-174	55	12.7	-	55	12.7	-	478	32.2	-	14.2	8.4	41
SL	174-175	102	26.5	-	102	26.5	-	-	-	-	-	-	-
Total(SL)		516	24	26.5	516	26.5	24	4526	62	52	146.6	89.1	39
SP	2-5	10000*	81	-	10000*	74	-	10000*	74	-	1hr	35min	41.7
SP	5-6	10000*	74	-	10000*	74	-	10000*	74	-	1hr	32min	47
SP	9-13	10000*	81	-	10000*	81	-	10000*	81	-	1hr	27min	55
SP	37-39	6	0	-	3699	77	-	3699	77	-	26min	22min	15
SP	39-40	2	0	-	-	-	-	-	-	-	-	-	-
SP	50-62	6188	82.9		6188	82.9	-	6188	82.9	-	35 min	21	40
SP	45-47	-	-	-	-	-	-	-	-	-	-	-	-
SP	47-50	2	0	-	2	0	-	-	-	-	-	-	-
SP	62-124	6	0	-	10000*	28	-	10000*	28	-	1hr*	58min	2
SP	40-45	-	-	-	-	-	-	-	-	-	-	-	-
Total(SL)		36206	79	74	49889	68	77	49889	68	77	5hr	3.25hr	35

Table 3: Time and irrelevant branches

S	$T_{static}(s)$	B_{E+I}	B_P	B_{Tot}
replace	4.5	90	57	181
siena	4.1	34	16	185
STPG	35	16	10	272
SL	0.4	33	11	383
SP	703	49	21	447

62% fewer runs (median 52%) and 39% fewer amount of time to find behavioral differences.

Neither Pex+eXpress nor Pex was able to find behavioral differences between some version pairs of class `StructParser` in 5 minutes (a bound that we use in our experiments for all subjects). For these version pairs, we increased the bound to 1 hour (or 10000 runs). Pex was not able to find behavioral differences for 4 version pairs even in 1 hour, while Pex+eXpress found behavioral differences for all these version pairs. If Pex was unable to detect behavioral differences within the bound of 1 hour, we put the time in the column T_{pex} as 1 hour and the number of runs as 10000 (the bound on the number of runs) to calculate the total in the last row of Table 2. Similarly, for the columns E_{Pex} and I_{Pex}, we take the number of runs as 10000 if conditions E and I, respectively, are not satisfied within 10000 runs.

$Pex + eXpress$ takes non-trivial average time of 700 seconds to find irrelevant branches for the class `StructParser` due to a large number of method invocations. However, considering that most of the changes cannot be covered in 1 hour, the time taken to find irrelevant branches is substantially less.

Changes between two version pairs (40-45 and 40-47) could not be covered by either Pex nor $Pex+eXpress$. One of the changes (between version pairs 47-50) was a refactoring. For this version pair, program state was infected but no behavioral differences were detected by either Pex or $Pex + eXpress$.

In summary, $Pex + eXpress$ was able to detect behavioral differences for four of the version pairs that could not be detected by Pex. On average, Pex was able to find behavioral differences in 68% fewer runs (median 77%) and 35% fewer amount of time. The reduction in number of runs is substantially larger than reduction in amount of time due to non-trivial time taken by $eXpress$ in finding irrelevant branches.

Seeding program exploration with existing tests. Table 4 shows the results obtained by using the existing test suite to seed the program exploration. Column C shows the class name. Column V

Table 4: Results obtained by seeding existing test suite for structorian

C	V	N_{Pex}/T	$N_{p_{seed}}/T$	$N_{eXpress}/T$	$N_{e_{seed}}/T$
SP	2-5	10000/1hr*	10000/1hr*	2381/35min	181/17min
SP	37-39	3699/26m	60/1m	851/22m	47/11m
SP	39-40	10000/1hr*	304/2m	10000/1hr*	251/12m
SP	45-47	10000/1hr*	10000/1hr*	10000/1hr*	10000/1hr*
SP	47-50	10000/1hr*	81/1m	10000/1hr*	64/10m
SP	62-124	10000/1hr*	59/1m	7228/58m	41/10m
SL	169-174	478/1m	324/1m	34/1m	18/1m
SL	150-169	299/1m	37/1m	52/1m	29/1m
SL	9-139	2988/2m	69/1m	1002/1m	52/1m
Total		64476/6.5hr	20934/2hr3m	41568/5hr9m	10683/2hr3m

*If behavior differences are not detected, we take the number of runs as 10,000 (the maximum number of runs that we ran our experiments with)

shows the pair of version numbers. The next four columns show the number of runs and time taken by the four techniques: Pex, Pex with seeding, Pex+eXpress, and Pex+eXpress with seeding, respectively, for finding behavioral differences. Note that DSE runs required by our Incremental Exploration also includes the seeded test runs.

In Table 4, if all the changed blocks are not covered, we take the number of runs as 10,000 (the maximum number of runs that we ran our experiments with). For 9 of the version pairs of `structorian` (out of 16 that we used in our experiments), the existing test suite of `structorian` could not find behavioral differences. Therefore, we consider these 9 version pairs for our experiments. Pex could not find behavioral differences for 5 of the 9 version pairs in 10,000 runs. Seeding the program exploration with the existing test suite helps Pex in finding behavioral differences for 3 of these version pairs under test. Pex+eXpress could not find behavioral differences for 3 of the 9 version pairs in 10,000 runs. Seeding the program exploration with the existing test suite helps Pex+eXpress in finding behavioral differences for 2 of these version pairs under test.

In summary, Pex requires around 67.5% of the original runs and 67% less time (required by Pex without test seeding) and Pex+eXpress requires around 74% of the original runs and 70% less time (required by Pex+eXpress without test seeding). In terms of time, Pex with seeding marginally wins over Pex+eXpress with seeding due to time taken by Pex+eXpress in finding irrelevant branches.

7. RELATED WORK

Previous approaches [10, 29, 15] generate regression unit tests

achieving high structural coverage on both versions of the class under test. However, these approaches explore all the irrelevant paths, which cannot help in achieving any of the conditions I or E in the PIE model [33]. In contrast, we have developed a new search strategy for DSE to avoid exploring these irrelevant paths.

Santelices et al. [2, 27] use data and control dependence information along with state information gathered through symbolic execution, and provide guidelines for testers to augment an existing regression test suite. Unlike our approach, their approach does not automatically generate tests but provides guidelines for testers to augment an existing test suite.

Some existing search strategies [4, 36] guide DSE to efficiently achieve high structural coverage in a program under test. However, these techniques do not specifically target covering a changed region. In contrast, our approach guides DSE to avoid exploring paths that cannot help in satisfying any of the conditions P, I, or E of the PIE model.

Differential symbolic execution [22] determines behavioral differences between two versions of a method (or a program) by comparing their symbolic summaries [11]. Summaries can be computed only for methods amenable to symbolic execution. However, summaries cannot be computed for methods whose behavior is defined in external libraries not amenable to symbolic execution. Our approach still works in practice when these external library methods are present since our approach does not require summaries. In addition, both approaches can be combined using demand-driven-computed summaries [1], which we plan to investigate in future work.

Li [19] prioritizes source code portions for testing based on dominator analysis. In particular, her approach finds a minimal set of blocks in the program source code, which, if executed, would ensure the execution of all of the blocks in the program. Howritz [14] prioritizes portions of source code for testing based on control and flow dependencies. These two approaches focus on testing in general. In contrast, our approach focuses specifically on regression testing.

Ren et al. develop a change impact analysis tool called Chianti [23]. Chianti uses a test suite to produce an execution trace for two versions of a program, and then categorizes and decomposes the changes between two versions of a program into different atomic types. Chianti uses only an existing test suite and does not generate new tests for regression testing. In contrast, our approach focuses on regression test generation.

Some existing capture and replay techniques [9, 21, 25] capture the inputs and outputs of the unit under test during system-test execution. These techniques then replay the captured inputs for the unit as less expensive unit tests, and can also check the outputs of the unit against the captured outputs. However, the existing system tests do not necessarily exercise the changed behavior of the program under test. In contrast, our approach generates new tests for regression testing.

Joshi et al. [16] use the path constraints of the paths followed by the tests in an existing test suite to generate inputs that violate the assertions in the test suite. The generated test inputs follow the same paths already covered by the existing test suite and do not explore any new paths. In contrast, our approach exploits the existing test suite to explore new paths. Majumdar and Sen [20] propose the concept of hybrid concolic testing. Hybrid concolic testing seeds the program with random inputs so that the program exploration does not get stuck at a particular program location. In contrast, our approach exploits the existing test suite to seed the program exploration. Since the existing test suite is expected to achieve a higher structural coverage, the existing test suite is expected to discover more hard-to-discover branching nodes in comparison with random inputs. Godefroid et al. [13] propose a DSE based approach for fuzz testing of large applications. Their approach uses a single seed for program exploration. In contrast, our approach seeds multiple tests to program exploration. Seeding multiple tests can help program exploration in covering the changes more efficiently as discussed in Section 5.5.2.

Law and Rothermel [18] propose an impact analysis technique, called PathImpact. PathImpact uses method execution traces to find out impacted methods when a method is modified. Our technique of building an inter-procedural graph by pruning certain method call chains that are found to be reachable to a changed region is similar to the technique. However, our technique works on a static interprocedural graph due to which our technique is safe in contrast to PathImpact. In addition, our technique further reduces the size of Inter-procedural graph by not adding the already visited methods from the inter-procedural graph.

Rothermel and Harrold [24] introduce the notion of dangerous edges and use these dangerous edges for regression test selection. Our irrelevant branches for execution and infection (set B_{E+I}) of a changed region is the inverse of these dangerous edges. However, our approach also finds irrelevant branches that cannot help in propagating a state infection to observable output.

Xu and Rothermel [38] propose a directed test generation technique that uses the existing test suite to cover parts of the program that are not covered by the existing test suite. In particular, the approach first collects the set of branches B that are not covered by the existing test suite. To cover a branch $b_i = < v_i, v_j > \in B$, the approach selects all the tests T that cover the vertex v_i. For each test $t_i \in T$, the approach collects the path constraints p_i of path followed by t_i until v_i, negates the predicate at v_i from p_i to get path condition p'_i. The approach then generates a test that covers the branch b_i by solving the path conditions p_i. However if all the path conditions of paths followed by the tests T are not solvable, the approach cannot generate a test to cover the branch b_i, which can furthermore compromise the coverage of additional branches. In contrast, our incremental exploration technique can still generate a test to cover such branches. In addition, the approach focuses only on satisfying condition E of the PIE model, while our approach helps in satisfying E, I, and P of the PIE model to find behavioral differences.

8. DISCUSSION

In this section, we discuss some of issues of the current implementation of our approach and how they can be addressed.

Added/Deleted and Refactored Methods. If a method M (or a field F) is added or deleted from the original program version, `eXpress` does not detect M (or F) as a changed region. The change is detected if a method call site (or reference to F) is added or deleted from the original program version. If the added or deleted method (or field) is never invoked, the behavior of the two versions is the same unless M is an overriding method. We plan to incorporate support for handling such overriding methods that are added or deleted. Similarly, if a method M is refactored between the two versions, `eXpress` does not detect M as a changed region. However, when a method is refactored, its call sites are changed accordingly (unless the method undergoes `Pull Up` or `Push Down` refactoring). Hence, `eXpress` detects the method containing call sites of M as changed. In our experiments, we considered versions of `replace` in which a method signature was changed, and versions of `structorian` in which a method was renamed.

Granularity of Changed Region. In our current implementation, a changed region is the list of continuous instructions that include

all the changed instructions in a method. One method can have only a single changed region. Hence, a changed region can be as big as a method and as small as a single instruction. The granularity of a changed region can be increased to a single method or reduced to single instruction. Changing the granularity to single method M can affect the efficiency of our approach in reducing DSE runs since some of the branches in M that should be considered irrelevant would not be considered irrelevant. In contrast, reducing the granularity to a single instruction makes our approach more efficient in reducing DSE runs. However, the overhead cost of our approach is increased due to state checking at multiple points in the program. In future work, we plan to enhance `eXpress` to allow users to choose from different levels of granularity.

Original/New Program Version. In our current implementation, we perform DSE on the new version of a program. We then execute a test (generated after each run) on the original version. We can also perform DSE on the original version instead of the new version. One approach may be efficient than the other depending on the types of changes made to the program. In future work, we plan to conduct experiments to compare the efficiency of the two approaches with respect to the types of changes.

Pruning of Branches for Propagation. In future work, we plan to prune more categories of branches that cannot help in Propagation (P). Consider that a changed region is executed and the program state is infected after the execution of the changed region; however, the infection is not propagated to any observable output. Let χ be the last location in the execution path such that the program state is infected before the execution of χ but not infected after its execution. χ can be determined by comparing the value spectra [35] obtained by executing the test on both versions of the program. This category contains all the branching nodes after the execution of χ. These branches can be obtained by inspecting the path P followed in the previous DSE run. Let $P = < b_1, b_2, .., b_\chi...b_n >$, where bi is the a branching node in Path P, while b_χ is the last branching node containing χ. We flip the branching nodes starting from b_χ to b_n in P until a branching node where the infected program state is propagated.

9. CONCLUSION

Regression testing aims at generating tests that detect behavioral differences between two versions of a program. To expose behavioral differences, a test execution needs to satisfy the conditions: Execution (E), Infection (I), and Propagation (P), as stated in the PIE model [33]. Dynamic symbolic execution (DSE) can be used to generate tests for satisfying these conditions. DSE explores paths in the program to achieve high structural coverage, and exploration of all these paths can often be expensive. However, many of these paths in the program cannot help in satisfying the three conditions in any way. In this paper, we presented an approach and its implementation called `eXpress` for regression test generation using DSE. `eXpress` prunes paths or branches that cannot help in detecting the E, I, or P condition such that these conditions are more likely to be satisfied earlier in path exploration. In addition, our approach can exploit the existing test suite for the original version to efficiently execute the changed regions (if not already covered by the test suite). Experimental results on various versions of programs showed that our approach can efficiently satisfy E, I, or P conditions than without using our approach.

10. REFERENCES

[1] S. Anand, P. Godefroid, and N. Tillmann. Demand-driven compositional symbolic execution. In *Proc. TACAS*, pages 367–381, 2008.

[2] T. Apiwattanapong, R. Santelices, P. K. Chittimalli, A. Orso, and M. J. Harrold. Matrix: Maintenance-oriented testing requirement identifier and examiner. In *Proc. TAICPART*, pages 137–146, 2006.

[3] C. Boyapati, S. Khurshid, and D. Marinov. Korat: Automated testing based on java predicates. In *Proc. ISSTA*, pages 123–133, 2002.

[4] J. Burnim and K. Sen. Heuristics for scalable dynamic test generation. In *Proc. ASE*, pages 443–446, 2008.

[5] C. Cadar, V. Ganesh, P. M. Pawlowski, D. L. Dill, and D. R. Engler. Exe: automatically generating inputs of death. In *Proc. CCS*, pages 322–335, 2006.

[6] L. Clarke. A System to Generate Test Data and Symbolically Execute Programs. *IEEE Trans. Softw. Eng.*, 2(3):215–222, 1976.

[7] D. Dig, C. Comertoglu, D. Marinov, and R. Johnson. Automatic detection of refactorings in evolving components. In *Proc. ECOOP*, pages 404–428, 2006.

[8] H. Do, S. G. Elbaum, and G. Rothermel. Supporting controlled experimentation with testing techniques: An infrastructure and its potential impact. *ESE*, pages 405–435, 2005.

[9] S. Elbaum, H. N. Chin, M. B. Dwyer, and M. Jorde. Carving and replaying differential unit test cases from system test cases. *TSE*, 35(1):29–45, 2009.

[10] R. B. Evans and A. Savoia. Differential testing: a new approach to change detection. In *Proc. FSE*, pages 549–552, 2007.

[11] P. Godefroid. Compositional dynamic test generation. In *Proc. POPL*, pages 47–54, 2007.

[12] P. Godefroid, N. Klarlund, and K. Sen. DART: directed automated random testing. *Proc. PLDI*, pages 213–223, 2005.

[13] P. Godefroid, M. Y. Levin, and D. A. Molnar. In *Proc. NDSS*, pages 151–166, 2008.

[14] S. Horwitz. Tool support for improving test coverage. In *Proc. ESOP*, pages 162–177, 2002.

[15] W. Jin, A. Orso, and T. Xie. Automated behavioral regression testing. In *Proc. ICST*, 2010.

[16] P. Joshi, K. Sen, and M. Shlimovich. Predictive testing: Amplifying the effectiveness of software testing. In *Proc. ESEC-FSE*, pages 561–564, 2007.

[17] J. C. King. Symbolic Execution and Program Testing. *Communications of the ACM*, 19(7):385–394, 1976.

[18] J. Law and G. Rothermel. Whole program path-based dynamic impact analysis. In *Proc. ICSE*, pages 308–318, 2003.

[19] J. J. Li. Prioritize code for testing to improve code coverage of complex software. In *Proc. ISSRE*, pages 75–84, 2005.

[20] R. Majumdar and K. Sen. Hybrid concolic testing. In *Proc. ICSE*, pages 416–426, 2007.

[21] A. Orso and B. Kennedy. Selective capture and replay of program executions. In *Proc.WODA*, pages 1–7, 2005.

[22] S. Person, M. B. Dwyer, S. Elbaum, and C. S. Păsăreanu. Differential symbolic execution. In *Proc. FSE*, pages 226–237, 2008.

[23] X. Ren, B. G. Ryder, M. Stoerzer, and F. Tip. Chianti: A change impact analysis tool for java programs. In *Proc. ICSE*, pages 664–665, 2005.

[24] G. Rothermel and M. J. Harrold. A safe, efficient regression test selection technique. *ACM Trans. Softw. Eng. Methodol.*, 6(2):173–210, 1997.

[25] D. Saff, S. Artzi, J. H. Perkins, and M. D. Ernst. Automatic test factoring for Java. In *Proc. ASE*, pages 114–123, 2005.

[26] R. A. Santelices, P. Chittimalli, T. Apiwattanapong, A. Orso, and M. J. Harrold. Test-suite augmentation for evolving software. In *Proc. ASE*, pages 218–227, 2008.

[27] R. A. Santelices, P. Chittimalli, T. Apiwattanapong, A. Orso, and M. J. Harrold. Test-suite augmentation for evolving software. In *Proc. ASE*, pages 218–227, 2008.

[28] K. Sen, D. Marinov, and G. Agha. CUTE: a concolic unit testing engine for C. In *Proc. FSE*, pages 263–272, 2005.

[29] K. Taneja and T. Xie. DiffGen: Automated regression unit-test generation. In *Proc. ASE*, pages 407–410, 2008.

[30] K. Taneja, T. Xie, N. Tillmann, J. de Halleux, and W. Schulte. Guided path exploration for regression test generation. In *Proc. ICSE, NIER*, May 2009.

[31] N. Tillmann and J. de Halleux. Pex-white box test generation for .NET. In *Proc. TAP*, pages 134–153, 2008.

[32] N. Tillmann and W. Schulte. Parameterized unit tests. In *Proc. ESEC/FSE*, pages 253–262, 2005.

[33] J. Voas. PIE: A dynamic failure-based technique. *TSE*, 18(8):717–727, 1992.

[34] T. Xie, D. Marinov, and D. Notkin. Rostra: A framework for detecting redundant object-oriented unit tests. In *Proc. ASE*, pages 196–205, 2004.

[35] T. Xie and D. Notkin. Checking inside the black box: Regression testing by comparing value spectra. *TSE*, 31(10):869–883, 2005.

[36] T. Xie, N. Tillmann, P. de Halleux, and W. Schulte. Fitness-guided path exploration in dynamic symbolic execution. In *Proc. DSN*, 2009.

[37] B. Xin, W. N. Sumner, and X. Zhang. Efficient program execution indexing. In *Proc. PLDI*, pages 238–248, 2008.

[38] Z. Xu and G. Rothermel. Directed test suite augmentation. In *APSEC*, pages 406–413, 2009.