

eXpress: Guided Path Exploration for Regression Test Generation

Kunal Taneja¹, Tao Xie¹, Nikolai Tillmann², Jonathan de Halleux², Wolfram Schulte²

¹Department of Computer Science, North Carolina State University, Raleigh, NC, 27695, USA

²Microsoft Research, One Microsoft Way, Redmond, WA, 98074, USA

¹{ktaneja, txie}@ncsu.edu, ²{nikolait, jhalleux, schulte}@microsoft.com

ABSTRACT

Regression test generation aims at generating a test suite that can detect behavioral differences between the original and the new versions of a program. Regression test generation can be automated by using Dynamic Symbolic Execution (DSE), a state-of-the-art test generation technique, to generate a test suite achieving high structural coverage. DSE explores paths in the program to achieve high structural coverage but exploration of all these paths can often be expensive. However, if our aim is to detect behavioral differences between two versions of a program, we do not need to explore all these paths in the program since not all these paths are relevant for detecting behavioral differences. In this paper, we propose an approach on guided path exploration that avoids exploring irrelevant paths in terms of detecting behavioral differences. Hence, behavioral differences are more likely to be detected earlier in path exploration. In addition, our approach leverages the existing test suite (if available) for the original version to efficiently execute the changed regions of the program and infect program state. Experimental results on ** versions of four programs show that our approach requires about **% fewer runs (i.e., explored paths) on average to cause the execution of a changed region and **% fewer to cause program-state differences after its execution than exploration without guidance. In addition, our approach requires **% fewer runs to cover all the changed regions (i.e., infection) by exploiting an existing test suite than exploration without using the test suite.

1. INTRODUCTION

Regression test generation aims at generating a test suite that can detect behavioral differences between the original and the new versions of a program. A behavioral difference between two versions of a program can be reflected by the difference between the observable outputs produced by the execution of the same test (referred to as a difference-exposing test) on the two versions. Developers can inspect these behavioral differences to determine whether they are intended or unintended (i.e., regression faults).

Regression test generation can be automated by using Dynamic Symbolic Execution (DSE) [11, 21, 4], a state-of-the-art test generation technique, to generate a test suite achieving high structural

coverage. DSE explores paths in a program to achieve high structural coverage, and exploration of all these paths can often be expensive. However, if our aim is to detect behavioral differences between two versions of a program, we do not need to explore all these paths in the program since not all these paths are relevant for detecting behavioral differences.

To formally investigate irrelevant paths for exposing behavioral differences, we adopt the Propagation, Infection, and Execution (PIE) model [26] of error propagation. According to the PIE model, a fault can be detected by a test if a faulty statement is executed (E), the execution of the faulty statement infects the state (I), and the infected state (i.e., error) propagates to an observable output (P). A change in the new version of a program can be treated as a fault and then the PIE model is applicable for effect propagation of the change. Many paths in a program often cannot help in satisfying any of the conditions P, I, or E of the PIE model.

In this paper, we present an approach¹ *eXpress* and its implementations that uses DSE to detect behavioral differences based on the notion of the PIE model.

Our approach first determines all the branches (in the program under test) that cannot help in achieving any of the conditions E and I of the PIE model in terms of the changes in the program. To make test generation efficient, we develop a new search strategy for DSE to avoid exploring these irrelevant branches (including which can lead to an irrelevant path²). In particular, our approach guides DSE to avoid from flipping branching nodes³, which on flipping execute some irrelevant branch.

In addition, our approach can exploit the existing test suite (if available) for the original version by seeding the tests in the test suite to the program exploration. Our technique of seeding the exploration with the existing test suite can be used to efficiently augment an existing test suite so that various changed parts of the program (that are not covered by the existing test suite) are covered by the augmented test suite.

This paper makes the following major contributions:

Path Exploration for Regression Test Generation. We propose

¹An earlier version of this work [23] is described in a four-page paper that appears in the NIER track of ICSE 2009. This version significantly extends the previous work in the following major ways. First, in this paper, we develop techniques for efficiently finding irrelevant branches that cannot execute any change. Second, we develop techniques for exploiting the existing test suite for efficiently generating regression tests. Third, we automate our approach by developing a tool. Fourth, we conduct extensive experiments to evaluate our approach.

²An irrelevant path is a path that cannot help in achieving P, I, and E of the PIE model.

³A branching node in the execution tree of a program is an instance of a conditional statement in the source code. A branching node consists of two sides (or more than two sides for a `switch` statement): the true branch and the false branch. Flipping a branching node is flipping the execution of the program from the true (or false) branch of the branching node to the false (or true) branch. Flipping a branching node representing a switch statement is flipping the execution of the current branch to another unexplored branch.

an approach called `eXpress` that uses DSE for efficient generation of regression unit tests. To the best of our knowledge, ours is the first approach that guides path exploration on the new version specifically for regression test generation.

Incremental Test Generation. We develop a technique for exploiting an existing test suite, so that path exploration focuses on covering the changes rather than starting from scratch. To the best of our knowledge, ours is the first technique that leverages an existing test suite for automated regression test generation.

Implementation. We have implemented our `eXpress` approach in a tool as an extension for Pex [24], an automated structural testing tool for .NET developed at Microsoft Research. Pex has been previously used internally at Microsoft to test core components of the .NET architecture and has found serious bugs [24]. The current Pex has been downloaded for thousands of times in industry. `eXpress` efficiently generates regression tests given an original and new version of a program.

Evaluation. We have conducted experiments on ** versions of four programs. Experimental results show that our approach requires about **% fewer runs (i.e., explored paths) on average to cause the execution of a changed region and **% fewer runs on average to cause program-state differences after its execution than exploration without guidance. In addition, our approach requires **% fewer runs to cover all the changed regions by exploiting an existing test suite than exploration without using the test suite.

2. DYNAMIC SYMBOLIC EXECUTION

In our approach, we use Pex [24] as a DSE-based tool. Pex starts program exploration with some default inputs. Pex then collects constraints on program inputs from the predicates at the branching statements executed in the program. We refer to these constraints at branching statements as branch conditions. The conjunction of all branch conditions in the path followed during execution of an input is referred to as a path constraint. Pex keeps track of the previous executions to build a dynamic execution tree. Pex, in the next run, chooses one of the unexplored branch in the execution tree (dynamically discovered thus far). Pex flips the chosen branching node in the dynamic execution tree to generate a new input that follows a new execution path. Pex uses various heuristics [29] for choosing a branching node (to flip next) in the execution tree using various search strategies with an objective of achieving high code coverage fast.

3. EXAMPLE

In this section, we illustrate the `eXpress` approach with an example. `eXpress` takes as input two versions of a program and produces as output a regression test suite. The test suite on execution detects behavioral differences (if any exist) between the two versions of program under test. Although `eXpress` analyzes assembly code of C# programs, in this section, we illustrate the `eXpress` approach using program source code.

Consider the example in Figure 1. Suppose that the statements at Lines 11, and 12 of `testMe` have been added in the new version. We next describe our approach using the example.

Difference Finder. The `Difference Finder` component of `eXpress` compares the original and the new versions of the program under test to find differences between each corresponding method of the two program versions. For the program in Figure 1, `Difference Finder` detects that the statements at Lines 11, and 12 are added in the new version.

Graph Builder. The `Graph Builder` component of `eXpress` then builds a Control-Flow Graph (CFG) of the new version of the

```

1  static public int TestMe(string[] c){
2      int state = 0;
3      for(int i=0; i< c.Length; i++){
4          if(c[i] == '['){
5              state =1;
6          }
7          else if(state == 1 && c[i] == "("){
8              state =2;
9          }
10         else if(state == 2 && c[i] == "("){
11             state =3;
12         }
13         else if(state == 3 && c[i] == "<"){
14             state =4;
15         }
16         else if(state == 4 && c[i] == ">"){ state =4;
17             //Added in new version
18             break; //Added in new version
19         }
20         if(c[i]==' '){
21             break;
22         }
23         if(state==1){
24             return state;
25         }
26     }
27     return state;
28 }

```

Figure 1: An example program

program under test and marks the changed vertices in the graph. Left side of Figure 2 shows the CFG of the example program in Figure 1. The labels of vertices in the CFG show the corresponding line numbers in Figure 1. The red (dark) vertices show the added statements at Lines 11 and 12. The gray vertices are the branching nodes (i.e. these vertices denote the branching statements in the program in Figure 1), while the white vertices show the other statements in the program.

Graph Traverser. The `Graph Traverser` component traverses the CFG to find each branch⁴ b in a program such that if b is taken, the program execution cannot reach a red (dark) vertex at Lines 11 or 12. On traversing the CFG in Figure 2, the `Graph Traversal` detects that taking the branches $\langle 2, 18 \rangle$ ⁵, $\langle 14, 15 \rangle$, $\langle 16, 17 \rangle$, and $\langle 12, 14 \rangle$ (dotted/dark red edges in Figure 2), the program execution cannot reach either of the vertices at Lines 11 or 12. Since after taking these branches, the execution cannot reach the changed statements at Line 11 and 12, the execution of these branches cannot help in executing the changed statements; behavioral differences between two versions of the program cannot be detected without executing the changed statements. Hence, these branches are not explored by the `Dynamic Test Generator` (described later in this section) component of `eXpress` while generating regression tests for the program under test.

Instrumenter. The `Instrumenter` component transforms the two versions of the program code such that the transformed program code is amenable to regression testing. In particular, the `Instrumenter` component instruments both versions of the program under test so that program states can be compared (to determine state infection) after the execution of the added statements at Lines 11 and 12.

Dynamic Test Generator. To cover the changed statement at Line 12, DSE needs at least 5 loop iterations. However, the number of iterations depend on the choice of branches DSE flips in each iteration. In each iteration DSE has the choice of flipping 7 branches in the program. For the program in Figure 1, Pex takes 441 DSE runs to cover the added statement at Line 12 (to infect the program state). The problem can be amplified if the number of states in the program increase. To reduce the branch search space of DSE, the `Dynamic Test Generator` component adopts the PIE model [26] of error propagation described in Section 1. In particular, the `Dynamic Test Generator` prunes all the following branches from the search space of DSE.

- **Branches not satisfying E.** The branches $\langle 2, 18 \rangle$, \langle

⁴A branch is an outgoing edge of a branching node

⁵ $\langle x, y \rangle$ represents a branch from node x to node y

Algorithm 1 *InterProceduralCFG(τ)*

Input: A test method τ .

Output: The inter-procedural Control Flow Graph (CFG) of the program under test.

```
1:  $Graph\ g \leftarrow GenerateIntraProceduralCFG(\tau)$ 
2: for all Vertex  $v \in g.Vertices$  do
3:   if  $v.Instruction = MethodInvocation$  then
4:      $c \leftarrow getMethod(v.Instruction)$ 
5:     if  $c \in MethodCallStack$  then
6:       goto Line 2 To avoid loops
7:     end if
8:     if  $c \in ReachableToChangedRegion$  then
9:        $g \leftarrow GraphUnion(ChangedMethod, g, v)$ 
10:      goto Line 2
11:     end if
12:     if  $c \in Visited$  then
13:       goto Line 2
14:     end if
15:     if  $c \in ChangedMethods$  then
16:        $ChangedMethod \leftarrow c$ 
17:       for all Method  $m \in MethodCallStack$  do
18:          $ReachableToChangedRegion.Add(m)$ 
19:       end for
20:     end if
21:      $MethodCallStack.Add(c)$ 
22:      $cg \leftarrow InterProceduralCFG(c)$ 
23:      $MethodCallStack.Remove(c)$ 
24:      $Visited.Add(c)$ 
25:      $g \leftarrow GraphUnion(cg, g, v)$ 
26:   end if
27: end for
28: return  $g$ 
```

PUT τ . An intra-procedural CFG g is constructed for the method τ . For each method invocation vertex⁹ (invoking Method c) in g , the algorithm *InterProceduralCFG* is invoked recursively with the invoked method c as the argument (Line 22 of Algorithm 1), while adding c to the call stack (Line 21). After the control returns from the recursive call, the method c is removed from the call stack (Line 23) and added to the set of visited methods (Line 24). The inter-procedural graph cg (with c as an entry method) resulting from the recursive call at Line 22 is merged with the graph g (Line 25). The algorithm *InterProceduralCFG* is not invoked recursively with c as argument in the following situations:

c is in call stack. If c is already in the call stack, *InterProceduralCFG* is not recursively invoked with c as argument (Lines 5-6). This technique ensures that our approach is not stuck in a loop in method invocations. For example, if method A invokes method B and method B invokes A. Then the construction of inter-procedural graph stops after A is encountered the second time.

c is already visited. If c is already visited, *InterProceduralCFG* is not recursively invoked with c as argument (Lines 23). This technique ensures that we do not have to build the same subgraph again.

c is in *ReachableToChangedRegion*. The set *ReachableToChangedRegion* is populated whenever a changed method¹⁰ is encountered. In particular, if a changed method is encountered, the methods currently in the call stack are added to the set *ReachableToChangedRegion* (Lines 17-19). If c is in *ReachableToChangedRegion*, *InterProceduralCFG* is not recursively invoked with c as argument, while merging CFG of some changed method with g (Line 8-11). Since our aim of building the intra-procedural CFG is to find irrelevant branches, those in the graph via which the execution cannot reach any changed instruction, the preceding three techniques help in achieving the aim while reducing the cost of building the inter-procedural CFG. In addition, the size of the inter-procedural CFG is reduced resulting in reduction in the cost of finding irrelevant branches.

4.3 Graph Traverser

The *Graph Traverser* component takes as input the inter-procedural CFG g constructed by the *Graph Builder* component and a set V

⁹A method invocation vertex is a vertex representing a call instruction.

¹⁰A changed method M_i is a method for which the set $\Delta_i \neq \phi$.

of changed vertices in the CFG g . *Graph Traverser* traverses the graph to find a set of branches B via which the execution cannot reach any of the branch in V . A branch b in CFG g is an edge $e = \langle v_i, v_j \rangle$: $e \in g$, where v_i is a vertex in g with a degree of more than one. The vertex v is referred to as a branching node. The pseudo code for finding the set of branches B is shown in Algorithm 2.

Algorithm 2 *FindUnreachableBranches(g, V)*

Input: A Graph g and a set V of vertices in Graph g .

Output: A set of branches B in Graph g that do not have a path to any vertex $v \in V$.

```
1:  $Reachable \leftarrow Reachable \cup V$ 
2: for all Vertex  $v \in g.Vertices$  such that  $v.degree > 1$  do
3:   if  $Reachable.Contains(v)$  then
4:     goto Line 2
5:   end if
6:    $\rho \leftarrow FindPathUsingDFS(v, V, Reachable)$ 
7:   if  $\rho$  is empty then
8:     for all Vertex  $n \in v.OutVertices$  do
9:        $B \leftarrow B \cup \langle v, n \rangle$ 
10:    end for
11:    goto Line 2
12:   end if
13:   for all Vertex  $pv \in \rho$  such that  $pv.degree > 1$  do
14:      $Reachable \leftarrow Reachable \cup pv$ 
15:   end for
16:   while  $v.OutVertices \neq \emptyset$  do
17:      $R \leftarrow R \cup \langle v, v_j \rangle$ :  $v_j \in \rho$ 
18:      $g \leftarrow g.RemoveEdge(v, v_j)$ 
19:      $\rho \leftarrow FindPathUsingDFS(v, V, Reachable)$ 
20:     if  $\rho$  is empty then
21:       for all Vertex  $o \in v.OutVertices$  do
22:          $B \leftarrow B \cup \langle v, o \rangle$ 
23:       end for
24:        $g \leftarrow g.AddEdges(R)$ 
25:       goto Line 2
26:     end if
27:   end while
28:    $g \leftarrow g.AddEdges(R)$ 
29: end for
30: return  $B$ 
```

For each Vertex v in g such that $degree(v) > 1$, the *Graph Traverser* performs a depth first search (DFS) from Vertex v (Line 6) to find a path between v and some vertex $c \in V$. If no path is found, all branches $b_i = \langle v, v_i \rangle$ are added to the set B (Lines 7-10) since none of these branches have a path to any of the vertices in V . If a path ρ is discovered (Lines 16-27 in Algorithm 2) by the DFS, there may still be a branch $b_i = \langle v, v_i \rangle$ that is not reachable to any vertex $c \in V$. To find such branch, we remove the edge $e_j = \langle v, v_j \rangle$: $e_j \in \rho$ from the graph and perform DFS again starting from v . We repeat the preceding steps until we either find no path ρ or all the edges from v are removed from the graph. If no path is found, the branches from v containing the remaining vertices are added to the set B (Lines 21-23).

To make the traversal efficient, whenever a path ρ is found, all the vertices r : $degree(r) > 1$ are added to the set of *Reachable*. This technique can help in making the future runs of DFS efficient. Whenever a vertex in the set *Reachable* is encountered, the DFS is stopped; returning the current path.

4.4 Instrumenter

The *Instrumenter* component uses Sets Δ_i (differences between method M_{i1} and M_{i2}) produced by the *Difference Finder*. For each changed method pair $\langle M_{i1}, M_{i2} \rangle$ for which $\Delta_i \neq \phi$, the *Instrumenter* component finds a region δ_i containing all the changed instructions in the program. δ_i is a minimal list of continuous instructions such that all the changed instructions in the method M_i are in the region δ_i . Hence, there can be a maximum of one changed region in one method. At the end of each changed region δ_i , the *Instrumenter* component inserts instructions to save the program state. In particular, the *Instrumenter* inserts the corresponding instructions for *PexStore* statements for each variable (and field) defined in the changed region. The *PexStore* statement for a variable x results in an assertion statement `PexAssert.IsTrue("uniqueName", x == currentX)` in the generated test, where

currentX is the value of x in the new version. The `Dynamic Test Generator` component generates tests for the new version $v2$. Once a test is generated that executes a changed region, the test is executed on Version $v1$ to compare program states after the execution of the changed region with the ones captured in the execution of Version $v2$. The instrumentation enables us to perform only one instance of DSE on the new version instead of performing two instances of DSE: one on the original and the other on the new program version. Performing two instances of DSE can be technically challenging since we have to perform the two DSE instances in a controlled manner such that both versions are executed with the same input and the execution trace is monitored for both the versions by a common exploration strategy to decide which branching node to flip next in the two versions.

4.5 Dynamic Test Generator

The `Dynamic Test Generator` component performs Dynamic Symbolic Execution (DSE) [5, 14, 11, 21, 4] to generate regression tests for the two given versions of a program. DSE iteratively generates test inputs to cover various feasible paths in the program under test (the new version in our approach). In particular, DSE flips some branching node discovered in previous executions to generate a test input for covering a new path. The node to be flipped is decided by a search strategy such as depth-first search. The exploration is quite expensive since there are an exponential number of paths with respect to the number of branches in a program. However, the execution of many branches often cannot help in detecting behavioral differences. In other words, covering these branches does not help in satisfying any of the condition E or I in the PIE model described in Section 1. Therefore, we do not flip such branching nodes in our new search strategy for generating test inputs that detect behavioral differences between the two given versions of a program. Recall that, we refer to such branches as irrelevant branches. These branches are found using the `Graph Traversal` component. We next describe the two categories of paths that our approach avoids exploring, and then describe their corresponding branches.

4.5.1 Paths being Pruned

Our approach avoids exploring the following categories of paths:

Category E. Let $V = v_1, v_2, \dots, v_n$ be the set of all vertices in CFG g such that $v_i \in g$ and $v_i.degree > 1$. Let $E_i = e_{i1}, e_{i2}, \dots, e_{in}$ be the set of outgoing edges (branches) from v_i . Let C be the set of changed vertices in g , $\rho(v_i, C, e_{ij})$ denotes a path (if exists) between a vertex v_i to any vertex in set C that takes the branch e_{ij} of Vertex v_i . For all vertices $v_i \in V$, our approach avoids from exploring all the branches $e_{ij} \in E_i : \rho(v_i, C, e_{ij}) = \phi$ (such branches are found using `Graph Traversal` component). In other words, if an instance of branching node v_i is in the discovered dynamic execution tree¹¹, such that the branch e_{ij} is not explored yet, our approach does not flip v_i to make program execution take branch e_{ij} . As an effect, our approach helps reduce the number of paths to be explored by avoiding the paths $P \in g : e_{ij} \in P$.

Category I. Let $C \in g$ be the set of nodes in all the changed regions. Let some $c_i \in C$ be executed i.e at least one instance c_{ij} of c_i is present in the dynamic execution tree T . Consider that the program state is not infected after the execution of changed region containing c_i . Our approach prunes all the branching nodes $B \in T$, such that for all $b_i \in B : \rho(c_{ij}, b_i) \neq \phi, b_i \notin C$, where $\rho(c_i, b_i)$ is the path between the vertices c_i and b_i (if exists) in T . As an

effect of pruning B , our approach does not explore the subpaths after the execution of a changed region that does not cause any state infection.

4.6 Incremental Test Generation

A regression test suite achieving high code coverage may be available along with the original version of a program. This test suite may be manually written or generated by an automated test generation tool for the original version. However, the existing test suite might not be able to cover all the changed parts of the new version of the program. Our approach can reuse the existing test suite so that changed parts of the program can be executed efficiently due to which test generation is likely to find behavior differences earlier in path exploration. Our approach executes the existing test suite to build an execution tree for the tests in the test suite. Our approach then starts the program exploration using the dynamic execution tree instead of starting from an empty tree. Our approach of seeding test inputs can help efficiently cover the changed parts of the program with two major reasons:

Discovery of hard to discover branching nodes. By seeding the existing test suite to Pex, our approach executes the test suite to build an execution tree of the program (discovered by executing the existing test suite). Some of the branches in the discovered execution tree may take a large number of DSE runs (without seeding any tests) to get discovered. Flipping some of these discovered branching nodes nearer in CFG to the changed parts of the program has more likelihood of covering the changed regions of the program [3]. Although, currently, our approach does not specifically prioritize the flipping of the branches near to the changed parts, our approach can help these branches to get discovered, which might take large number of DSE runs as shown in the example in Section 3.

Priority of DSE to cover not covered parts of the program. DSE techniques employ branch prioritization so that a high coverage can be achieved faster due to which DSE techniques choose a branch from the execution tree (discovered thus far) that have a high likelihood of covering new program parts. By seeding the existing test suite to program exploration, the DSE techniques do not waste time on covering the parts of the program already covered by the existing test suite. Instead, the DSE techniques give priority to branching nodes that can cover new parts of the program, which include the changed parts. Hence, the changed parts are likely to be covered earlier in path exploration.

5. EXPERIMENTS

We conducted experiments on four programs and their 68 versions (in total) collected from three different sources. In our experiments, we try to address the following research questions:

RQ1. How much less number of DSE runs does Pex require to execute the changed regions between the two versions of a program with the assistance of `express`?

RQ2. How much less number of DSE runs does Pex require to infect the program states with the assistance of `express`?

RQ3. How much more number of tests does Pex, with the assistance of `express`, generate that execute changed regions between the two versions of a program?

RQ4. How much more number of tests does Pex, with the assistance of `express`, generate that infect the program states?

RQ5. How much less number of DSE runs does Pex require to cover the changed regions when the program exploration is seeded with the existing test suite?

5.1 Subjects

¹¹ A dynamic execution tree is the tree formed from the paths executed in the previous executions. Multiple instances $c_{i1}, c_{i2}, \dots, c_{in}$ of a node c_i in CFG can be present in a dynamic execution tree due to loops in a program.

To answer the research questions, we conducted experiments on four subjects. Table 1 shows the details about the subjects. Column 1 shows the subject name. Column 2 shows the number of classes in the subject. Column 3 shows the number of classes that are covered by tests generated in our experiments. Column 4 shows the number of versions (not including the original version) used in our experiments. Column 5 shows the number of lines of code in the subject.

`replace` and `siena` are programs available from the Subject Infrastructure Repository (SIR) [7]. `replace` and `siena` are written in *C* and *Java*, respectively. `replace` is a text processing program, while `siena` is an Internet-scale event notification program. We chose these two subjects (among the others available at the SIR) in our experiments as we could convert these subjects into C# using Java 2 CSharp Translator¹². We could not convert other subjects available at the SIR with the exception of `tcas`. The experimental results on `tcas` are presented in the previous version of this work [23]. We seeded all the 32 faults available for `replace` at the SIR one by one to generate 32 new versions of `replace`. For `siena`, SIR contains eight different sequentially released versions of `siena` (versions 1.8 through 1.15). Each version provides enhanced functionalities or corrections with respect to the preceding version. We use these eight versions in our experiments. In addition to these eight versions, there are nine seeded faults available at SIR. We seeded all the nine faults available at SIR one by one to synthesize nine new versions of `siena`. In total, we conduct experiments on these 17 versions of `siena`. For `replace`, we use the `main` method as a PUT for generating tests. For `siena`, we use the methods `encode` (for changes that are transitively reachable from `encode`) and `decode` (for changes that are transitively reachable from `decode`) in the class `SENP` as PUTs for generating tests. The method `encode` requires non-primitive arguments. Existing Pex cannot handle non-primitive types effectively but provides support for writing factory methods for non-primitive types. Hence, we manually wrote factory methods for the non-primitive types in `SENP`. In particular, we wrote factory methods for classes `SENP`, `Packet`, `Event`, and `Filter`. Each factory method invokes a sequence (of length up to three) of the state-modifying public methods in the corresponding class. The parameters for these methods, and the length of the sequence (up to three) are passed as inputs to the factory methods. During exploration, Pex generates concrete values for these inputs to cover various parts of the program under test.

STPG¹³ is an open source program hosted by the codeplex website, Microsoft’s open source project hosting website¹⁴. The codeplex website contains snapshots of check-ins in the code repositories for STPG. We collect three different versions of the subject STPG from the three most recent check-ins. We use the `Convert(string path)` method as the PUT for generating tests as `Convert` is the main conversion method that converts a string path data definition to a `PathGeometry` object.

`structorian`¹⁵ is an open source binary data viewing and reverse engineering tool. `structorian` is hosted by Google’s open source project hosting website¹⁶. The website also contains snapshots of check-ins in the code repositories for `structorian`. We collected all the versions of snapshots for the classes `StructLexer`,

Table 1: Experimental subjects

Project	Classes	Classes Covered	Versions	LOC
replace	1	1	32	625
STPG	1	1	2	684
siena	6	6	17	1529
structorian	70	8	21	6561

`BaseLexer` and `StructParser`. We chose these classes in our experiments due to three reasons. First, these classes have several revisions available in the repository. Second, these classes are of non trivial size and complexity. Third, these classes have corresponding tests available in the repository. For the classes `StructLexer` and `StructParser`, we generalized one of the available concrete test methods by promoting primitive types to arguments and removing the assertions. We used these generalized test methods as PUTs for our experiments. `structorian` contains a manually written test suite. We use this test suite for seeding the exploration for addressing the question RQ5.

For addressing questions RQ1-RQ4 we use all the four subjects, while for addressing the question RQ5 we use `structorian` because of two major reasons: first, `structorian` has a manually written test suite that can be used to seed the exploration. Second, revisions of `structorian` contains non trivial changes that cannot be covered by existing test suite. Hence, our approach of seeding the program exploration is useful for covering these changes. `replace` contains changes to one statement due to which most of the changes can be covered by the existing test suite. `siena` and STPG do not have an existing test suite for us to use.

5.2 Experimental Setup

For `replace` and `siena`, we find behavioral differences between the original version and each version $v2$ synthesized from the available faults in the SIR. We use `express` and the default search strategy in Pex [24, 29] to find behavioral differences. In addition to the versions synthesized by seeding faults, we also find behavioral differences between each successive versions of `siena` (versions 1.8 through 1.15) available in SIR, using `express` and the default search strategy in Pex [24, 29]. For STPG and `structorian`, we find behavioral differences between two successive pairs of versions that we collected.

To address RQ1, we compare the number of runs of DSE required by the default search strategy in Pex with the number of runs required by Pex enhanced with `express` to execute a changed region. To address RQ2, we compare the number of runs required by the default search strategy in Pex with the number of runs required by Pex enhanced with `express` to infect the program states. To address RQ3, we compare the number of tests that cover a changed region generated by Pex enhance with `express` with the number of such tests generated by default search strategy in Pex. If more number of tests are generated that cover a changed region, it is easier for developers (or testers) to debug the program under test (if the changes are faulty) and gives more confidence to developers that the changes they made do not introduce any unwanted side effects. To address RQ4, we compare the number of tests that infect the program state after the execution of changed region generated by Pex enhance with `express` with the number of such tests generated by default search strategy in Pex. To address RQ5, we compare the number of DSE runs required by the default search strategy in Pex (and `express`) to cover all the blocks in all the changed regions with and without seeding the program exploration (with the existing test suite).

Currently, we have not automated the steps to prune branches that cannot help in achieving I of the PIE model. To simulate the

¹²<http://sourceforge.net/projects/j2cstranslator/>

¹³<http://stringtopathgeometry.codeplex.com/>

¹⁴<http://www.codeplex.com>

¹⁵<http://code.google.com/p/structorian/>

¹⁶<http://code.google.com>

pruning of branches to achieve I, in our experiments, we manually instrument the new version to throw an exception immediately after the changed regions, if the program state is not infected after the execution of the changed region. If the changed region is located inside a loop, we throw the exception immediately after the loop. In future work, we plan to automate the pruning of branches that cannot help in satisfying I. The rest of the approach is fully automated and is implemented in a tool called `eXpress`. We developed `eXpress` as an extension¹⁷ to Pex [24]. We developed its components to statically find irrelevant branches as a .NET Reflector¹⁸ AddIn.

5.3 Experimental Results

Table 2 shows the experimental results. Due to space constraints, we only provide the total and average values for the subjects `replace`, `siena`, and `STPG`. The detailed results for experiments on all the versions of these subjects are available on our project web¹⁹. However, we provide detailed results for `structorian` in this paper.

Column *S* shows the name of the subject. For `structorian`, the column shows the class name. Column *V* shows the number of version pairs for which we conducted experiments for the subject. For `structorian`, the column shows the version numbers on which experiment was conducted. These version numbers are the revision numbers in the google code repository of `structorian`. Column *E_{Pex}* shows the total number of DSE runs required by the default search strategy in Pex for satisfying E. Column *E_{eXpress}* shows the total number of DSE runs required by `eXpress` for satisfying E. Column *E_{Red}* shows the percentage reduction in the number of DSE runs by `eXpress` for achieving E. Column *N_{ePex}* shows the total number of tests, that execute a changed region, generated by Pex. Column *N_{eXpress}* shows the total number of tests, that execute a changed region, generated by `eXpress`. Column *N_{eInc}* shows the percentage increase in the number of generated tests that execute a changed region. Column *I_{Pex}* shows the total number of DSE runs required by the default search strategy in Pex for satisfying I. Column *I_{eXpress}* shows the total number of DSE runs required by `eXpress` for satisfying I. Column *I_{Red}* shows the percentage reduction in the number of DSE runs by `eXpress` for achieving I. Column *N_{iPex}* shows the total number of tests, generated by Pex, that infect the program state. Column *N_{iXp}* shows the total number of tests, generated by `eXpress`, that infect the program state. Column *N_{iInc}* shows the percentage increase in the number of generated tests that execute a changed region.

Table 3 shows the time taken for finding the irrelevant branches, time taken to generate tests, and the number of irrelevant branches found. Column *S* shows the subject. Column *T_{static}* shows the average time taken by `eXpress` to find irrelevant branches that cannot help in satisfying E of the PIE model. Column *T_{Pex}(s)* shows the average time taken by Pex to generate tests. Column *T_{eXpress}* shows the average time taken by `eXpress` to generate tests. Column *B_{Irr}* shows the average number of irrelevant branches that cannot help in satisfying E of the PIE model. In general, irrelevant branches are more if changes are towards the beginning of the PUT since there are likely to be more branches in the program that do not have a path to any changed regions. These branches also include the branches whose branching condition is not dependent on the inputs of the program and therefore do not lead to branching conditions during path exploration. Hence, pruning these branches is not helpful in making the DSE efficient. Column *B_{Tot}* shows

the total number in the CFG.

Results of `replace`. For the `replace` subject, among the 32 pairs of versions, the changed regions cannot be executed for 4 of theses versions (Versions 14, 18, 27, and 31) by the default strategy in Pex or by `eXpress` in 1000 DSE runs. We do not include these versions while calculating the sum of DSE runs for satisfying I and E of the PIE model. For 3 of the versions (Versions 3, 22 and 32), the changed region was executed but the program state is not infected in 1000 DSE runs. We do not include these versions while calculating the sum of DSE runs for satisfying I of the PIE model). For 3 of the versions (Versions 12, 13, and 21), the changes are in the fields due to which there are no benefits of using `eXpress`. We exclude these three versions from the experimental results shown in Table 2.

`eXpress` takes around 5.7 seconds (on average) to find the irrelevant branches for each version of `replace` using optimizations. We also observe that the time varies for different versions (between 0.3 to 21.4 seconds) as our optimizations depend on the location of a change. In total, `eXpress` took 51.6% fewer runs in executing the changes with a maximum of 77.6% for Versions 23 and 24 available in the SIR. For these versions `eXpress` takes 95 DSE runs in contrast to 425 runs taken by Pex to execute the changed locations. In addition, `eXpress` took 46% fewer runs, in infecting the program state, with a maximum of 73.8% for Version 6 available in the SIR. For this version, `eXpress` takes 83 DSE runs in contrast to 317 runs taken by Pex to infect the program state after the execution of changed locations.

Results of `siena`. We observe that the changes in seven of the versions of `siena` are covered within ten runs by the default search strategy in Pex and `eXpress`. For these changes, there is no reduction in the number of runs. However, the number of generated tests that cover a changed region increase by a significant amount while using `eXpress` as compared to the default search strategy in Pex. The reason for the preceding phenomenon is that these changes are nearer to the entry vertex in the control flow graph. Hence, these changes can be covered in a relatively small number of runs. Moreover, for these types of changes, `eXpress` finds relatively large number of irrelevant branches because many of the branches in the CFG after these changes need not be explored to execute the changed region. As a result, test generation focuses on flipping significantly fewer branches (that are near to the change) due to which the tests that cover a changed region increase significantly. In two of the versions, changed regions were not covered by either `eXpress` or the default search strategy in Pex. An exception is thrown by the program before these changes could be executed. Pex and `eXpress` are unable to generate a test input to avoid the exception. Two of the changes are refactoring due to which the program state is never infected. In summary, `eXpress`, executed the changed region in 42% less runs to execute the changes as compared to the default search strategy in Pex and generates 121.1% more tests that execute the changed regions. In addition, `eXpress` infects the program state in 39.4% less runs and generates 170.2% more tests that infect the program state.

Results of `structorian`. The first six rows show the experimental results for changes in the class `StructLexer`. Rows 7 and 8 show the experiments for changes in the class `BaseLexer`, while the last 11 rows show the experimental results on versions of the class `StructParser`. For the versions of `StructLexer`, `eXpress` takes 24.6% less runs to execute a changed region than the default search strategy in Pex. In addition `eXpress` generates 43.3% more tests that cover a changed region than the default search strategy in Pex. In addition, `eXpress` infects the program state in 24.7% less runs and generates 39.7% more tests that infect the program state.

¹⁷<http://pex.codeplex.com/>

¹⁸<http://www.red-gate.com/products/reflector/>

¹⁹<https://sites.google.com/site/asergpr/projects/express/>

Table 2: Experimental Results

S	V	Execution						Infection					
		E_{Pex}	$E_{eXpress}$	$E_{Red}(\%)$	N_{ePex}	$N_{eXpress}$	$N_{eInc}(\%)$	I_{Pex}	$I_{eXpress}$	$I_{Red}(\%)$	Ni_{Pex}	$Ni_{eXpress}$	$Ni_{Inc}(\%)$
replace	32	1946	789	59.4	1183	2249	90	3203	1716	46.4	358	579	62
siena	17	286	166	42	549	1214	121.1	284	172	39.4	336	908	170.2
STPG	2	341	250	26.1	38	48	26.3	378	255	32.4	10	13	30
Total	51	2573	1205	53.1	1770	3511	98.4	3865	2143	44.6	704	1500	113.1
-structorian-													
SL	2-9	102	75	26.5	24	38	58.3	102	75	26.5	24	38	58.3
SL	9-139	102	75	26.5	24	38	58.3	152	107	29.6	8	11	37.5
SL	139-150	102	75	26.5	24	38	58.3	102	75	26.5	13	18	38.5
SL	150-169	53	46	13.2	20	25	25	53	46	13.2	20	25	25
SL	169-174	X	X	X	X	X	X	X	X	X	X	X	X
SL	174-175	102	75	26.5	24	38	58.3	-	-	-	-	-	-
SL	175-184	19	15	21.1	41	48	17.1	21	21	0	13	17	30.8
BL	45-174	2	2	0	999	999	0	3	3	0	243	265	9.1
BL	174-175	2	2	0	999	999	0	3	3	0	243	265	9.1
SP	2-5	-	1866	-	-	-	-	-	2587	-	-	-	-
SP	5-6	-	2587	-	-	-	-	-	2587	-	-	-	-
SP	9-13	-	1866	0	-	-	-	-	-	-	-	1866	-
SP	37-39	X	X	0	X	X	X	X	X	X	X	X	X
SP	39-40	X	X	0	X	X	X	X	X	X	X	X	X
SP	50-62	6188	1053	-	-	-	-	-	-	-	-	-	-
SP	45-47	2	2	0	43	53	23.3	2	2	0	43	53	23.3
SP	47-50	2	2	0	43	53	23.3	2	2	0	43	53	23.3
SP	62-124	2	2	0	43	53	23.3	2	2	0	43	53	23.3
SP	124-125	2	2	0	43	53	23.3	2	2	0	43	53	23.3
SP	125-166	-	7452	-	-	-	-	-	7452	-	-	-	-
SP	40-45	-	8214	-	-	-	-	-	8276	-	-	-	-

Table 3: Time and Irrelevant Branches

S	$T_{static}(s)$	T_{Pex}	T_{eXp}	B_{Irr}	B_{Tot}
replace	5.83	151.2s	139.5s	90	181
siena	4.11	78s	75s	34	185
STPG	35	176s	173s	16	272
structorian (SL)	0.47	135 s	131s	33	383
structorian (BL)	0.5	148 s	151s	9	75
structorian (SP)	703	1:27 hr	45 min	49	447

The changes in `BaseLexer` were just after the CFG entry vertex due to which all the generated tests execute the changed region. Both `eXpress` and default search strategy were not able to cover any changed region for six of the versions of class `StructParser` in 1000 DSE runs (a bound that we use in our experiments for all subjects). For these versions, we increased the bound to 10,000 runs. For five of these versions (4 and 5), default search strategy was not able to execute the changed region even in 10,000 runs, while `eXpress` executes the changed 4 regions and infect the program state for all of these versions. `eXpress` takes a non-trivial time of 700 seconds to find irrelevant branches for the class `StructParser` due to a large number of method invocations. However, considering that most of the changes cannot be covered even in 10,000 runs by `Pex` (more than 2.5 hours of exploration) the time taken to find irrelevant branches is significantly less.

Seeding program exploration with existing tests. Table 4 shows the results obtained by using the existing test suite to seed the program exploration. Column *C* shows the class name. Column *V* shows the pair of version numbers. The next four columns show the number of runs taken by the four techniques: `Pex`, `Pex` with seeding, `eXpress`, and `eXpress` with seeding, respectively, for covering all blocks in all changed regions. For nine of the version pairs of `structorian` (out of 21 that we used in our experiments), the existing test suite of `structorian` could not cover all blocks of the changed regions. Therefore, we consider these nine version pairs for our experiments. `Pex` could not cover all the blocks in the changed regions for six of these version pairs in 10,000 runs. Seeding the program exploration with the existing test suite, helps `Pex`

Table 4: Results obtained by seeding existing test suite for structorian

C	V	N_{Pex}	N_{Pseed}	$N_{eXpress}$	N_{eSeed}
SP	2-5	-	-	2381	181
SP	37-39	1355	60	851	47
SP	39-40	-	304	-	251
SP	45-47	-	-	-	-
SP	47-50	-	81	1341	64
SP	62-124	-	59	2067	41
SL	169-174	34	18	34	18
SL	150-169	53	37	46	29
SL	9-139	-	69	1089	52

in covering all the blocks in under 100 runs for four of these version pairs under test. There is a considerable reduction in the number of runs in the other version pairs with the exception of versions 2 – 5 and 45 – 47 in which seeding cannot help `Pex` in covering all the blocks in changed regions.

6. RELATED WORK

Previous approaches [9, 22] generate regression unit tests achieving high structural coverage on both versions of the class under test. However, these approaches explore all the irrelevant paths, which cannot help in achieving any of the conditions I, or E in the PIE model. In contrast, we have developed a new search strategy for DSE to avoid exploring these irrelevant paths.

Santelices et al. [2, 20] use data and control dependence information along with state information gathered through symbolic execution, and provide guidelines for testers to augment an existing regression test suite. Unlike our approach, their approach does not automatically generate tests but provides guidelines for testers to augment an existing test suite. Some existing search strategies [3, 29] guide DSE to efficiently achieve high structural coverage in a program under test. However, these techniques do not specifically target covering a changed region. In contrast, our approach guides DSE to avoid exploring paths that cannot help in executing a changed region. In addition, our approach avoids exploring paths that cannot help in P or I of the PIE model [26].

Differential symbolic execution [17] determines behavioral differences between two versions of a method (or a program) by com-

paring their symbolic summaries [10]. Summaries can be computed only for methods amenable to symbolic execution. However, summaries cannot be computed for methods whose behavior is defined in external libraries not amenable to symbolic execution. Our approach still works in practice when these external library methods are present as our approach does not require summaries. In addition, both approaches can be combined using demand-driven-computed summaries [1], which we plan to investigate in future work. Ren et al. develop a change impact analysis tool called Chianti [18]. Chianti uses a test suite to produce an execution trace for two versions of a program, and then categorizes and decomposes the changes between two versions of a program into different atomic types. Chianti uses only an existing test suite and does not aim to exercise behavioral differences between the two versions of the program under test. In contrast, the goal of our approach is specifically to expose behavioral differences across versions.

Some existing capture and replay techniques [8, 16, 19] capture the inputs and outputs of the unit under test during system-test execution. These techniques then replay the captured inputs for the unit as less expensive unit tests, and can also check the outputs of the unit against the captured outputs. However, the existing system tests do not necessarily exercise the changed behavior of the program under test. In contrast, our approach generates new tests to specifically find behavioral differences.

Joshi et al. [13] use the path conditions from the paths followed by the tests in an existing test suite to generate inputs that violate the assertions existing in the test suite. The generated test inputs follow the same paths already covered by the existing test suite and do not cover any new paths. In contrast, our approach exploits the existing test suite to cover new paths. Majumdar et al. [15] propose the concept of hybrid concolic testing. Hybrid concolic testing seeds the program with random inputs so that the program exploration does not get stuck at a particular program location. In contrast, our approach exploits the existing test suite to seed the program exploration. Since the existing test suite is expected to achieve a higher structural coverage, the existing test suite is expected to discover more hard to discover branching nodes in comparison with random inputs. Godefroid et al. [12] propose a DSE based approach for fuzz testing of large applications. Their approach uses a single seed for program exploration. In contrast, our approach seeds more than one tests to program exploration. Seeding multiple tests can help program exploration in covering the changes more efficiently as discussed in Section 4.6.

7. DISCUSSION

In this section, we discuss some of the limitations of the current implementation of our approach and how they can be addressed. **Added/Deleted and Refactored Methods.** If a method M (or a field F) is added or deleted from the original program version, `express` does not detect M (or F) as a changed region. The change is detected if a method call site (or reference to F) is added or deleted from the original program version. If the added or deleted method (or field) is never invoked, the behavior of the two versions is the same unless M is an overriding method. We plan to incorporate support for such overriding methods that are added or deleted. Similarly, if a method M is refactored between the two versions, `express` does not detect M as a changed region. However, when a method is refactored, its call sites are changed accordingly (unless the method undergoes Pull Up or Push Down refactoring). Hence, `express` detects the method containing call sites of M as changed. In our experiments, we considered versions of `replace` in which method signature was changed, and versions of `structorian` in which a method was renamed.

Granularity of Changed Region. In our current approach, a changed region is the list of continuous instructions that include all the changed instructions in a method. One method can have only a single changed region. Hence, a changed region can be as big as a method and as small as a single instruction. The granularity of a changed region can be increased to a single method or reduced to single instruction. Changing the granularity to single method M can affect the efficiency of our approach in reducing DSE runs since some of the branches in M that should be considered irrelevant would not be considered irrelevant. In contrast, reducing the granularity to a single instruction makes our approach more efficient in reducing DSE runs. However, the overhead cost of our approach is increased due to state checking at multiple points in the program. In future work, we plan to enhance `express` to allow users to choose from different levels of granularity.

Original/Modified Program Version. In our current approach, we perform DSE on the new version of a program. We then execute the test (generated after each run) on the new version. We can also perform DSE on the original version instead of the new version. One approach may be efficient than the other depending on the types of changes made to the program. In future work, we plan to conduct experiments to compare the efficiency of the two approaches with respect to the types of changes.

Branch Prioritization. `express` currently prunes branches that cannot help in detecting behavioral differences between the two versions. However, some branches in the program code can be more promising in detecting behavioral differences than others. Branching nodes can be prioritized based on the distance of the branching node to a changed region in the CFG. The distance $d(n1, n2)$ between any two nodes $n1$ and $n2$ in a CFG g is the number of nodes with degree > 1 between $n1$ and $n2$ in the shortest path between $n1$ and $n2$. Hence, the distance between a node b and a changed region Δ is the number of nodes with degree of more than one between b and the node representing the first instruction in Δ . The intuition behind this prioritization is that shorter the distance between the branching node and δ , it is likely to be easier to generate inputs to cause the execution of the changed region δ . This kind of branch prioritization is used by Burnim and Sen [3]. We can also prioritize branching nodes based on the probability to cause infection and to propagate the infection to an observable output. Moreover, we can prioritize branches based on data dependence from a changed region.

Pruning of Branches for Propagation. Currently, `express` prunes branches that cannot help satisfy E or I of the PIE model for change propagation. In future work, we plan to prune more categories of branches that cannot help in Propagation (P). Consider that a changed region is executed and the program state is infected after the execution of the changed region; however, the infection is not propagated to any observable output. Let χ be the last location in the execution path such that the program state is infected before the execution of χ but not infected after its execution. χ can be determined by comparing the value spectra [27, 28] obtained by executing the test on both versions of the program. This category contains all the branching nodes after the execution of χ . These branches can be obtained by inspecting the path P followed in the previous DSE run. Let $P = \langle b_1, b_2, \dots, b_\chi \dots b_n \rangle$, where b_i are the branching nodes in the Path P , while b_χ is the last instance of branching node containing χ . We do not flip the branching nodes from b_χ to b_n in P until if the program state is not propagated after the execution of χ .

8. CONCLUSION

Regression testing aims at generating tests that detect behavioral

differences between two versions of a software program. Dynamic symbolic execution (DSE) can be used to generate such difference exposing tests. DSE explores paths in the program to achieve high structural coverage, and exploration of all these paths can often be expensive. However, many of these paths in the program cannot help in detecting behavioral differences in any way. In this paper, we presented an approach and its implementation called `express` for regression test generation using DSE. `express` prunes paths or branches that cannot help in detecting behavioral differences such that behavioral differences are more likely to be detected earlier in path exploration. In addition, our approach can exploit the existing test suite for the original version to efficiently cover the changed regions (if not already covered by the test suite). Experimental results on various versions of programs showed that our approach can efficiently detect behavioral differences than without using our approach.

9. REFERENCES

- [1] S. Anand, P. Godefroid, and N. Tillmann. Demand-driven compositional symbolic execution. In *Proc. TACAS*, pages 367–381, 2008.
- [2] T. Apiwattanapong, R. Santelices, P. K. Chittimalli, A. Orso, and M. J. Harrold. Matrix: Maintenance-oriented testing requirement identifier and examiner. In *Proc. TAICPART*, pages 137–146, 2006.
- [3] J. Burnim and K. Sen. Heuristics for scalable dynamic test generation. In *Proc. ASE*, pages 443–446, 2008.
- [4] C. Cadar, V. Ganesh, P. M. Pawlowski, D. L. Dill, and D. R. Engler. Exe: automatically generating inputs of death. In *Proc. CCS*, pages 322–335, 2006.
- [5] L. Clarke. A System to Generate Test Data and Symbolically Execute Programs. *IEEE Trans. Softw. Eng.*, 2(3):215–222, 1976.
- [6] D. Dig, C. Comertoglu, D. Marinov, and R. Johnson. Automatic detection of refactorings in evolving components. In *Proc. ECOOP*, pages 404–428, 2006.
- [7] H. Do, S. G. Elbaum, and G. Rothermel. Supporting controlled experimentation with testing techniques: An infrastructure and its potential impact. *ESE*, pages 405–435, 2005.
- [8] S. Elbaum, H. N. Chin, M. B. Dwyer, and M. Jorde. Carving and replaying differential unit test cases from system test cases. *TSE*, 35(1):29–45, 2009.
- [9] R. B. Evans and A. Savoia. Differential testing: a new approach to change detection. In *Proc. FSE*, pages 549–552, 2007.
- [10] P. Godefroid. Compositional dynamic test generation. In *Proc. POPL*, pages 47–54, 2007.
- [11] P. Godefroid, N. Klarlund, and K. Sen. DART: directed automated random testing. *Proc. PLDI*, pages 213–223, 2005.
- [12] P. Godefroid, M. Y. Levin, and D. A. Molnar. In *NDSS*, 2008.
- [13] P. Joshi, K. Sen, and M. Shlimovich. Predictive testing: Amplifying the effectiveness of software testing. In *Proc. ESEC-FSE*, pages 561–564, 2007.
- [14] J. C. King. Symbolic Execution and Program Testing. *Communications of the ACM*, 19(7):385–394, 1976.
- [15] R. Majumdar and K. Sen. Hybrid concolic testing. In *Proc. ICSE*, pages 416–426, 2007.
- [16] A. Orso and B. Kennedy. Selective capture and replay of program executions. In *Proc. WODA*, pages 1–7, 2005.
- [17] S. Person, M. B. Dwyer, S. Elbaum, and C. S. Păsăreanu. Differential symbolic execution. In *Proc. FSE*, pages 226–237, 2008.
- [18] X. Ren, B. G. Ryder, M. Stoerzer, and F. Tip. Chianti: A change impact analysis tool for java programs. In *Proc. ICSE*, pages 664–665, 2005.
- [19] D. Saff, S. Artzi, J. H. Perkins, and M. D. Ernst. Automatic test factoring for Java. In *Proc. ASE*, pages 114–123, 2005.
- [20] R. A. Santelices, P. Chittimalli, T. Apiwattanapong, A. Orso, and M. J. Harrold. Test-suite augmentation for evolving software. In *Proc. ASE*, pages 218–227, 2008.
- [21] K. Sen, D. Marinov, and G. Agha. CUTE: a concolic unit testing engine for C. In *Proc. FSE*, pages 263–272, 2005.
- [22] K. Taneja and T. Xie. DiffGen: Automated regression unit-test generation. In *Proc. ASE*, pages 407–410, 2008.
- [23] K. Taneja, T. Xie, N. Tillmann, J. de Halleux, and W. Schulte. Guided path exploration for regression test generation. In *Proc. ICSE, NIER*, May 2009.
- [24] N. Tillmann and J. de Halleux. Pex-white box test generation for .NET. In *Proc. TAP*, pages 134–153, 2008.
- [25] N. Tillmann and W. Schulte. Parameterized unit tests. In *Proc. ESEC/FSE*, pages 253–262, 2005.
- [26] J. Voas. PIE: A dynamic failure-based technique. *TSE*, 18(8):717–727, 1992.
- [27] T. Xie and D. Notkin. Checking inside the black box: Regression testing based on value spectra differences. In *Proc. ICSM*, pages 28–37, 2004.
- [28] T. Xie and D. Notkin. Checking inside the black box: Regression testing by comparing value spectra. *TSE*, 31(10):869–883, 2005.
- [29] T. Xie, N. Tillmann, P. de Halleux, and W. Schulte. Fitness-guided path exploration in dynamic symbolic execution. In *Proc. DSN*, 2009.